

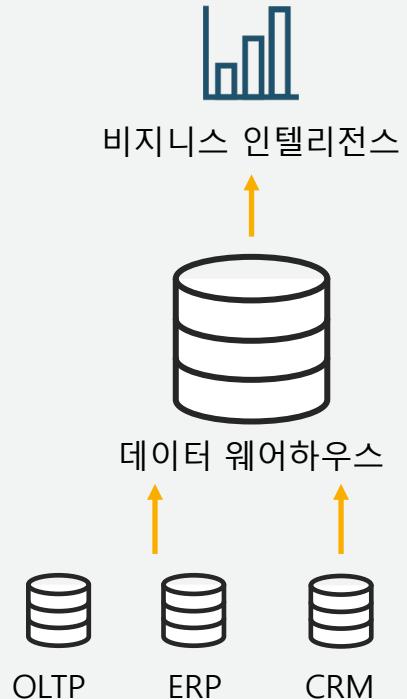


# Data Warehouse

Yoojeong Choi, Database SA, AWS

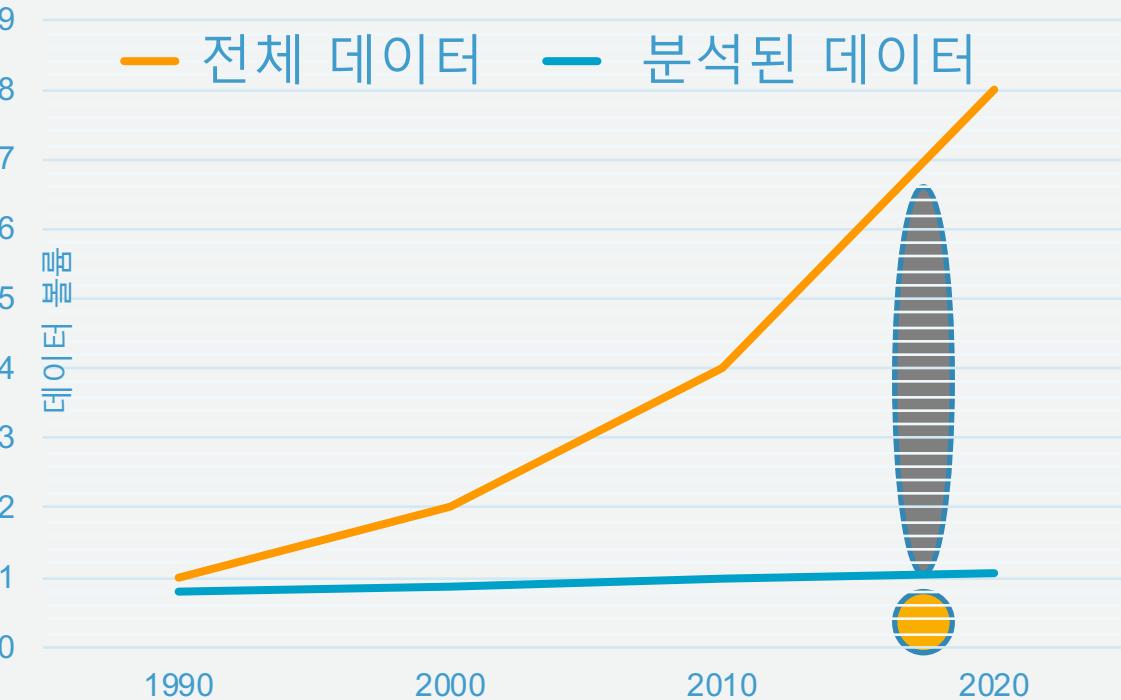


# 전통적인 데이터 Analytic 접근 방법



- 정형 데이터 형식
- 수TB ~ 수PB 바이트 크기
- 데이터 로딩 전 스키마 지정
- 지정된 인원이 한정된 자원을 기반으로 활용
- 높은 초기 투자 비용 및 유지보수 비용

# 초기 아키텍처로 인한 Dark Data 발생



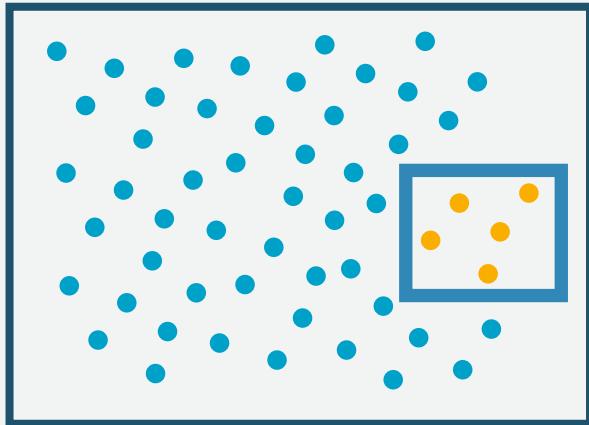
많은 데이터들이  
비즈니스에  
활용되지 못하고  
있습니다.

Sources:  
Gartner: User Survey Analysis: Key Trends Shaping the Future of Data Center Infrastructure Through 2011  
IDC: Worldwide Business Analytics Software 2012–2016 Forecast and 2011 Vendor Shares

# 페러다임의 전환

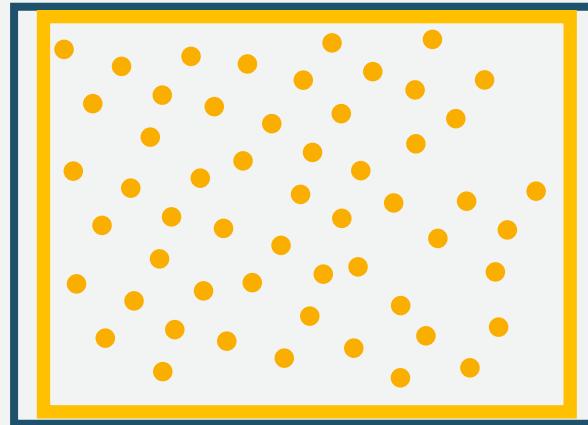
기업의 모든 데이터를 기반으로한 비니지스 인사이트 도출이 필요

전통적인 데이터 웨어하우스



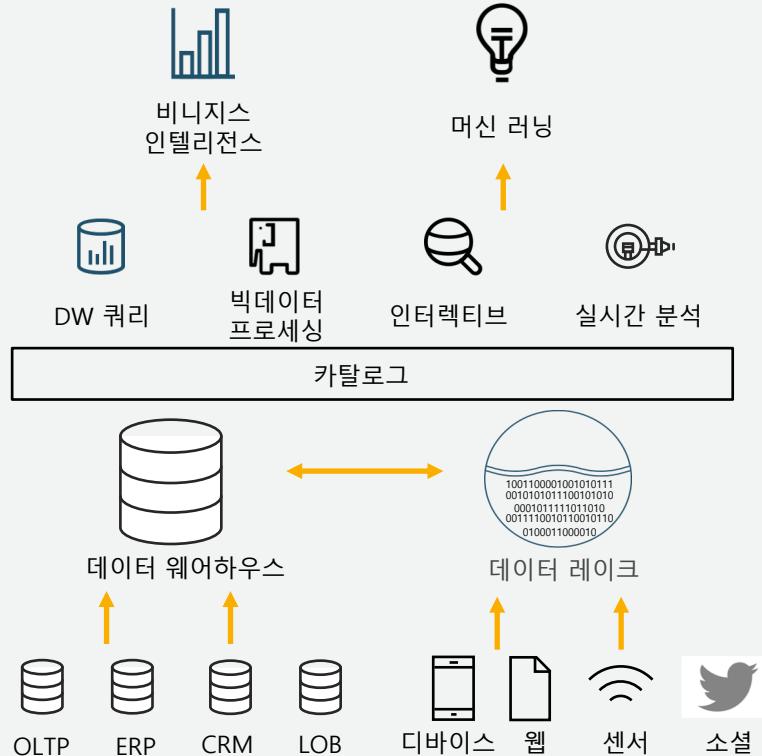
데이터 웨어하우스에  
저장된  
데이터만 분석

모던 데이터 웨어하우스



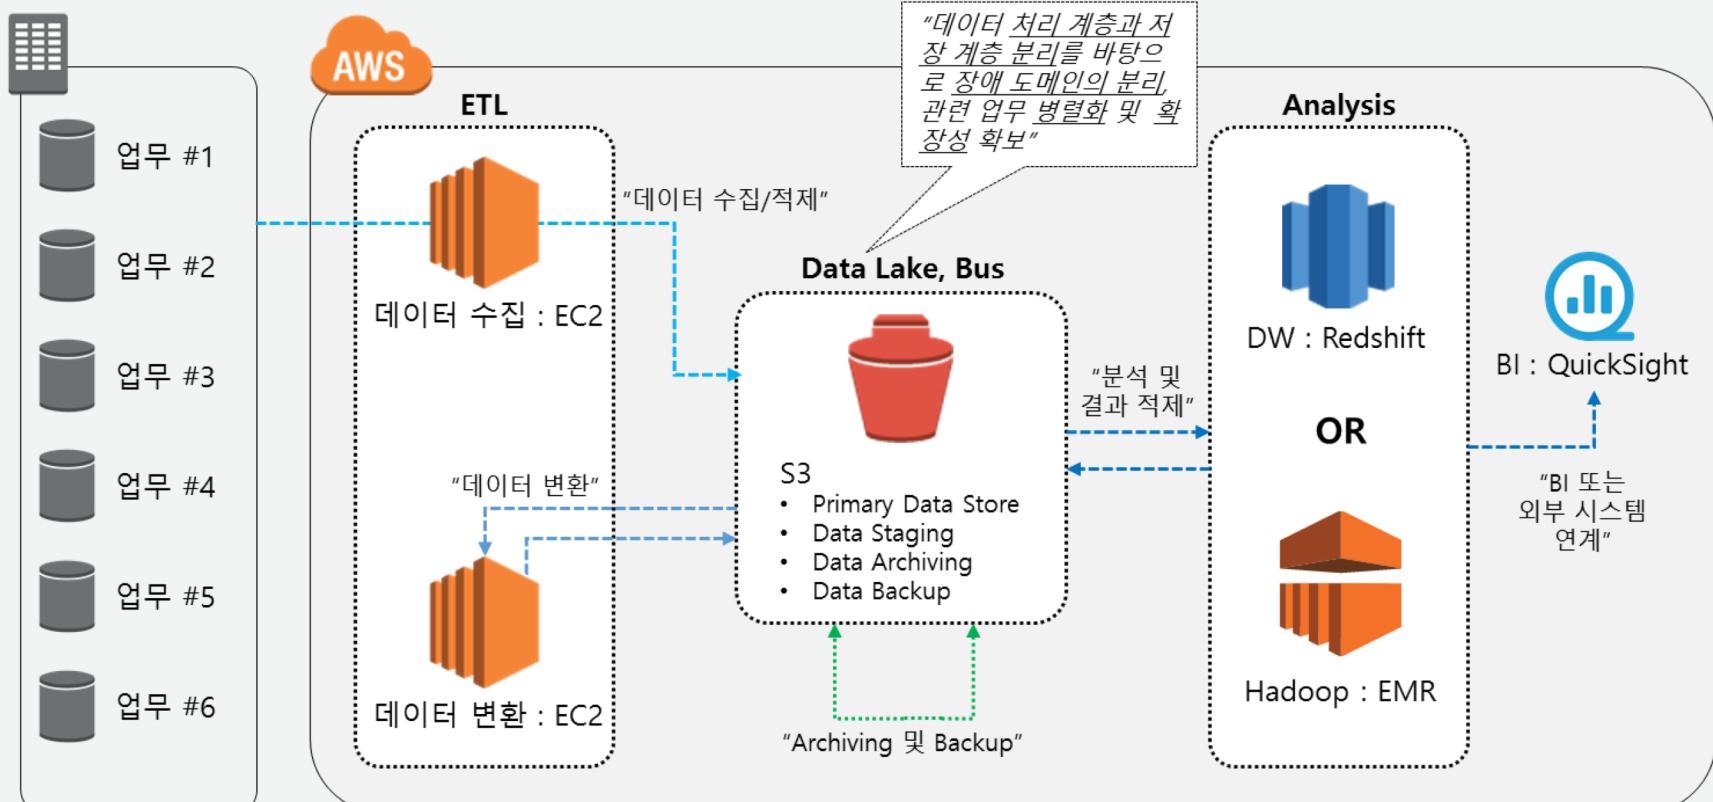
데이터레이크를 기반으로  
모든 데이터에 대한 분석

# 데이터 레이크를 기반한 데이터 분석 역량의 강화



- 정형, 비정형 데이터 분석
- 수TB ~ 수EB 바이트 크기
- 분석 시 스키마 지정
- 다양한 분석 도구
- 비용 최적화

# 만약 빅데이터 분석을 처음 시작하신다면...



# Amazon Redshift

1/10 비용으로 빠르고, 손쉬운 확장을 지원하는 완전 관리형 DW

## 빠른 성능



 "Hive 대비  
50~100배 빠른 성능"

 "210억  
행과 100억행  
조인에 2시간"

## 비용 효율적



"시간당 \$0.25로  
시작"

"년간 미압축  
데이터 기준 1TB에  
\$250~\$333"

## 보안



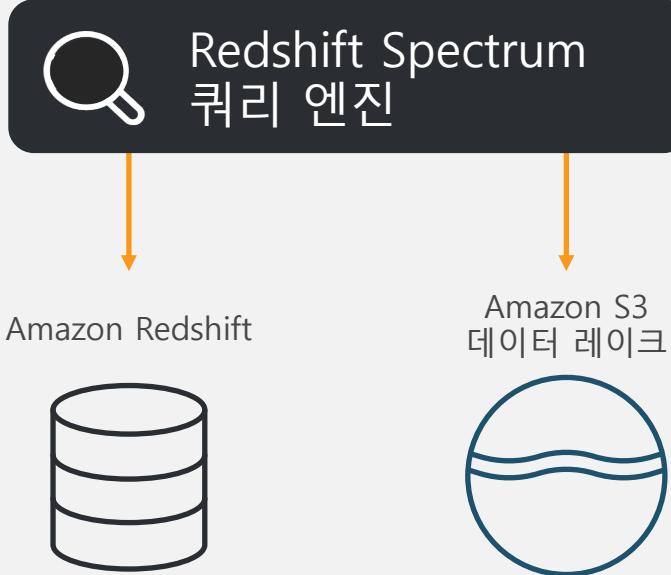
## 데이터 레이크와의 통합



"S3 (데이터레이크)  
대한 쿼리"

"AWS 다양한 서비스  
와의 유연한 통합"

# 데이터레이크와의 통합



- Amazon Redshift에서 S3데이터에 대한 쿼리
- Compute과 Storage 에 대한 별도의 용량 관리
- 빠른 쿼리 성능
- 무제한 병렬화
- CSV, ORC, Gro, Avro, Parquet 지원
- On demand, 쿼리당 스캔 데이터에 기준한 비용 (TB 당 \$5)

# 아키텍처

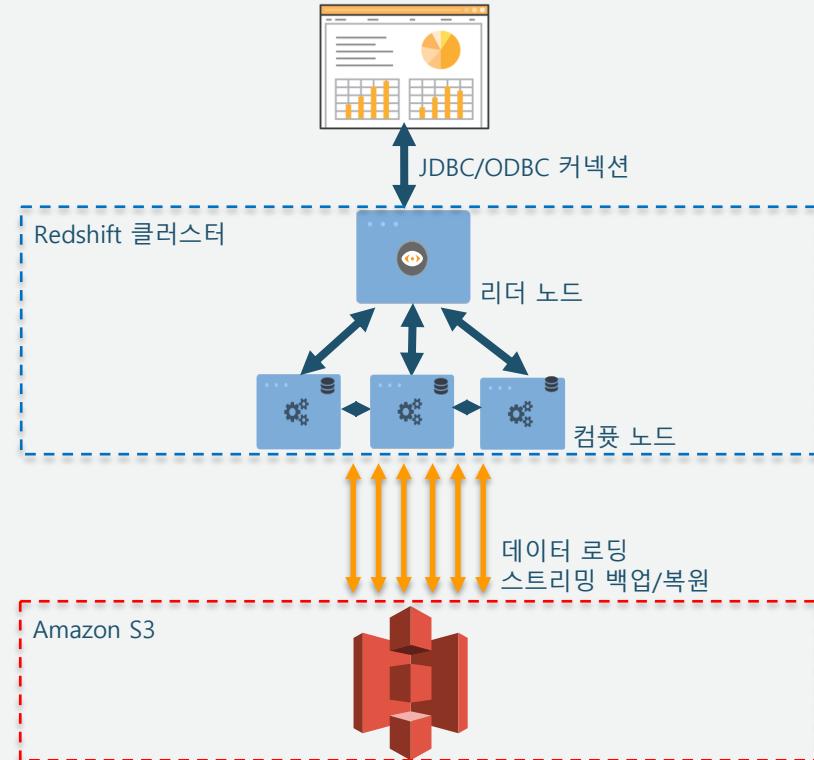
# Amazon Redshift 클러스터 아키텍처

MPP 기반의 Shared Nothing 아키텍처  
Amazon S3로의 지속적 엔 데이터 백업  
리더 노드

- SQL 엔드포인트
- 메타데이터 저장소
- 병렬 SQL 쿼리 코디네이터

컴퓨팅 노드

- 로컬, 컬럼 기반의 스토리지
- 병렬 쿼리 수행 주체
- 데이터 로드, 백업 및 복원 주체
- 2, 16, 32개 슬라이스



# 노드와 Slice



## 노드

### DC1 - Do Not Use

### DC2

- Based on EC2 i3 Instance Family
- SSD based storage
- Intel Broadwell Processors
- DC2.large
- DC2.8xlarge

### DS2

- Based on EC2 d2 Instance Family
- Magnetic based storage
- Intel Haswell Processors
- DS2.xlarge
- DS2.8xlarge



## 슬라이스

Redshift is a distributed system:

- A compute node contains slices (one per virtual core)
- A slice contains data

Slices are chosen based on two types of distribution:

- Round Robin (automated)
- Based on a distribution key (hash of a defined column)

Queries run on all slices in parallel: optimal query throughput can be achieved when data is evenly spread across slices

# 새로운 고밀도 컴퓨팅 노드 - DC2

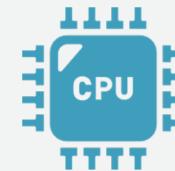
동일 가격에 DC1 대비 3배 높아진 I/O 성능과 30% 개선된 컴퓨팅 성능



NVMe SSD



DDR4 Memory



Intel E5-2686 v4 (*Broadwell*)



**Liberty Mutual.**  
INSURANCE

*"Amazon Redshift DC1에서 DC2로 이전한  
후 월간 레포트 작성이 9배가 빨라  
졌습니다."*

- Bradley Todd,  
Technical Architect, Liberty Mutual



# 결과 캐싱 - 일초 미만의 쿼리 응답 속도

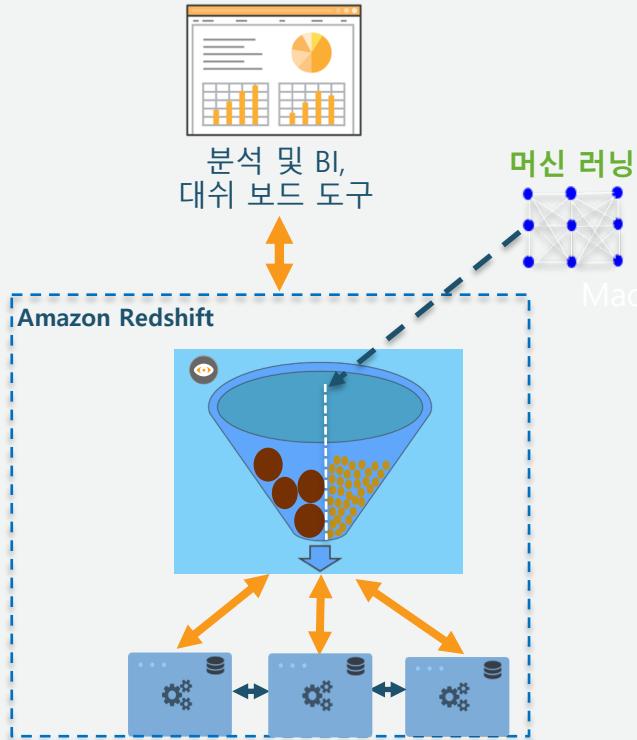
쿼리 실행 시간을 단축을 위한 쿼리 결과 메모리 캐싱



- 기본 활성화 (관련 파라메터 enable\_result\_cache\_for\_session)
- 캐시 조건
  - 쿼리에 사용되는 객체에 대한 액세스 권한 보유
  - Redshift 쿼리 관련 파라메터 및 조회하는 테이블내 데이터의 미변경
  - 쿼리 실행 시마다 실행해야 하는 함수 (예: GETDATE) 미사용
- 리더 노드 사이즈 및 쿼리 결과 크기를 기준으로 캐싱
- 세션에 관계 없이 캐쉬 조회
- SVL\_QLOG 뷰의 "source\_query" 항목을 통해 확인

# 단기 쿼리 가속화 Short Query Acceleration (SQA)

장기 쿼리로 인한 체감 성능 저하 완화 및 동시성 확보



- 머신 러닝 알고리즘을 기반으로 단기 쿼리 판단
  - 장기 쿼리뒤에서 단기 쿼리가 대기하지 않도록 SQA 전용 큐에서 단기 쿼리 실행
- SQA 조건
- GUI에서는 WLM의 총 슬롯 수가 15개 이하여야 함 (CLI, API를 통해 SQA를 활성화 한 경우 해당 없음)
  - CTAS 및 읽기 전용 쿼리에 사용
  - 기본값은 5초 (권장), 최대 20초까지 설정
- 'STL\_WLM\_Query'의 'service\_class' 컬럼값이 14인 쿼리가 SQA를 통해 Acceleration된 쿼리 문이며, 'final\_state'를 통해 'Completed' 또는 'Evicted' 여부를 확인할 수 있음

# Workload Management (WLM)

## Concurrent Query Execution

Long-running queries may cause short-running queries to wait in a queue  
As a result users may experience unpredictable performance

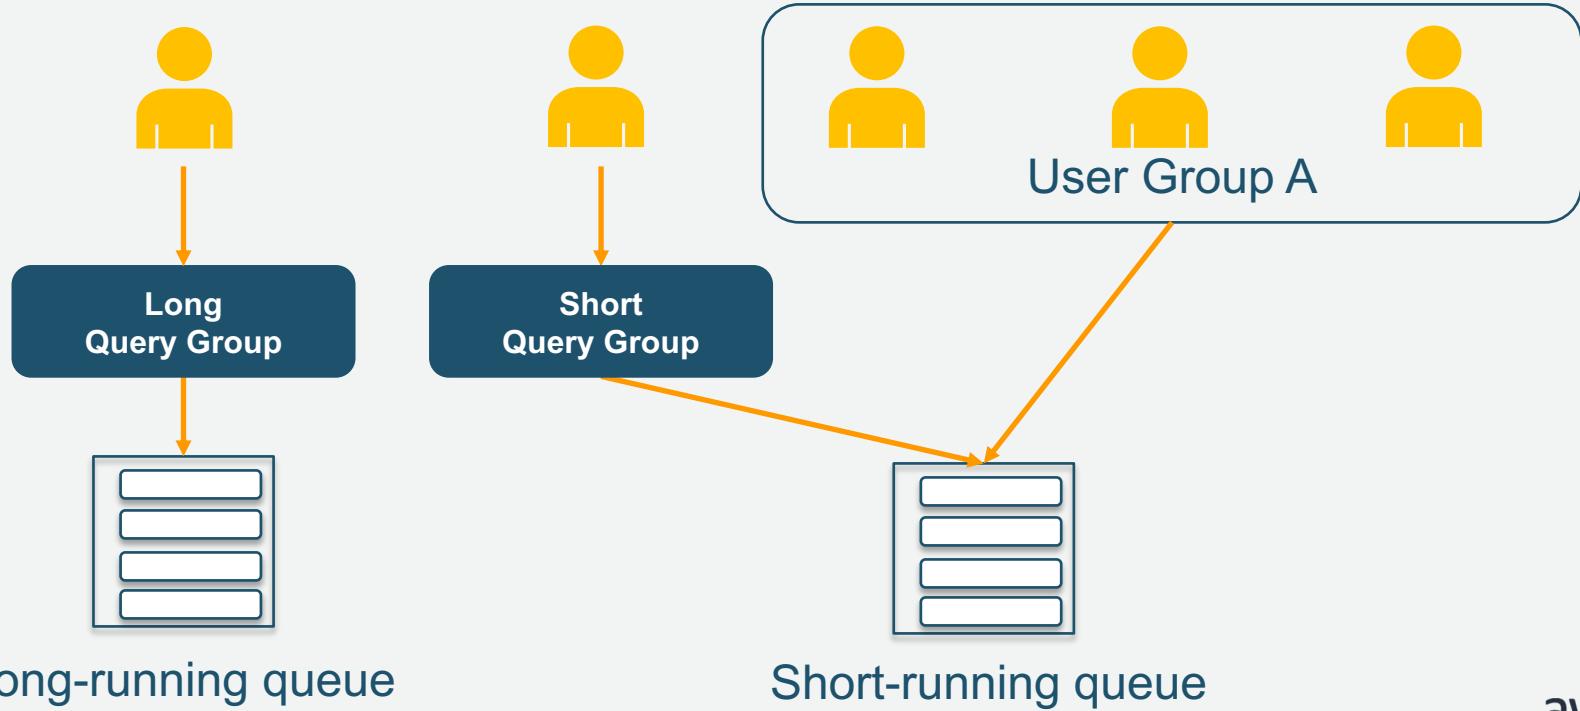
By default, a Redshift cluster is configured with one queue with 5 execution slots by default  
(50 max, 동시 연결 수량은 500)

전체 가용 메모리는 설정된 쿼리 Slot/동시성 를 기준으로 전체 메모리를 배분



# Workload Management (WLM)

실행 시간 및 비지니스 목적성에 따라 쿼리를 적합한 대기열로 라우팅



# Workload Management (WLM)

- 결과 캐싱 및 단기 쿼리 가속화는 WLM을 거치지 않음으로 성능 뿐 아니라 병렬성을 높히는데 큰 영향을 줌
- 병렬성을 높히면 그 만큼 단위 쿼리의 리소스가 줄어듬
- 메모리 부족 여부 확인  
'SVL\_QUERY\_SUMMARY' 또는 'v\_my\_last\_query\_summary'의 'is\_diskbased' 컬럼 확인

# 노드 선택의 기준

- IO Latency 요구 사항 및 필요 스토리지 용량
- 쿼리의 수행 패턴 (Slice Level, Node Level)
  - Slice Level vs Node Level
  - Broadcast
- 결과 캐싱 사용 여부
- 리더 노드 전용 함수 및 최종 Sorting, Aggregation (집계) 의 사용
- 최신 세대의 노드

## 리더 노드 전용 함수

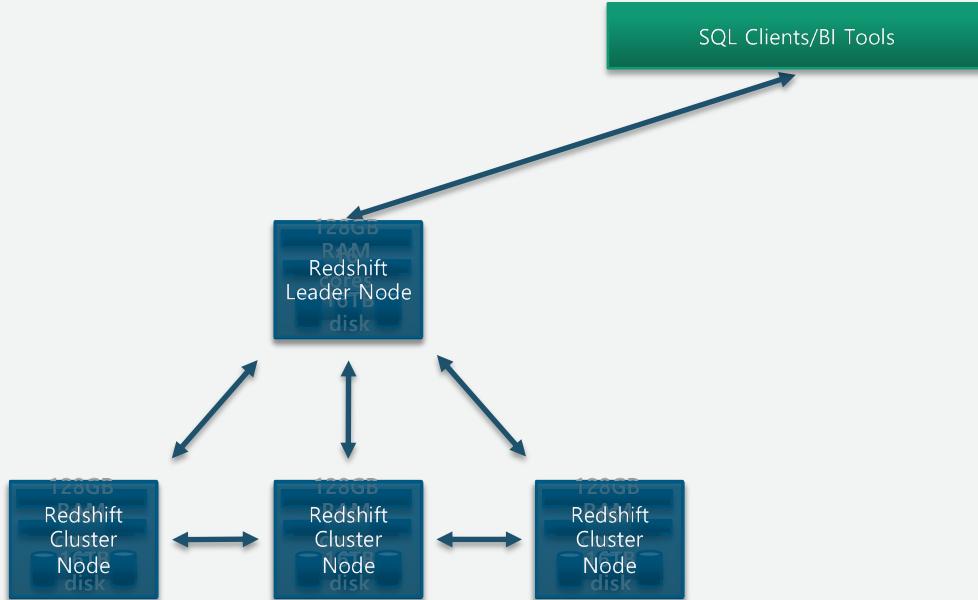
### 날짜 함수

- AGE
- CURRENT\_TIME
- CURRENT\_TIMESTAMP
- LOCALTIME
- ISFINITE
- NOW

### 문자열 함수

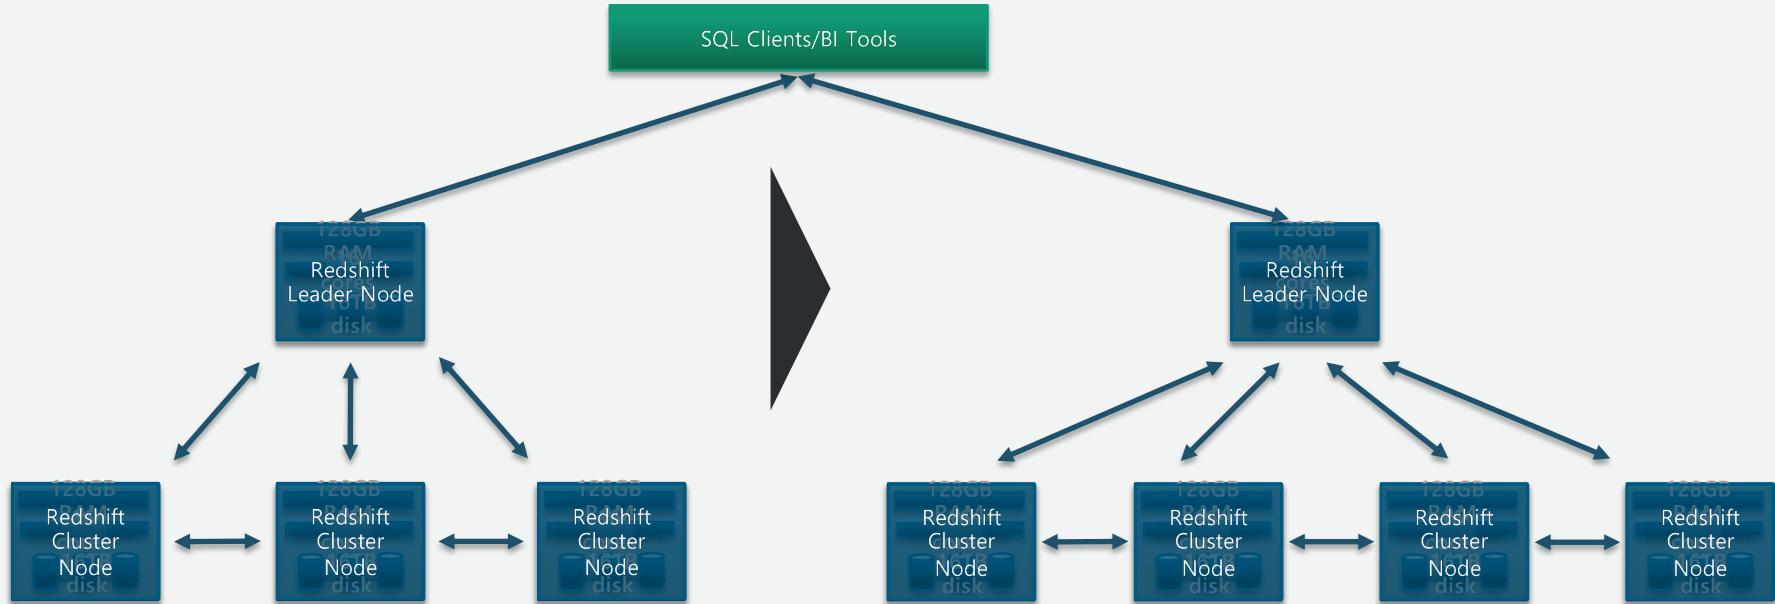
- ASCII
- GET\_BIT
- GET\_BYT
- SET\_BIT
- SET\_BYT
- TO\_ASCII

# 클러스터 및 노드 사이즈 변경



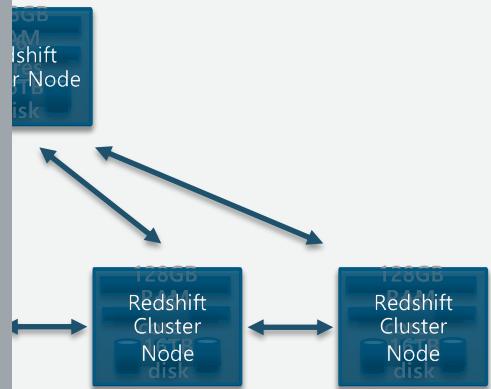
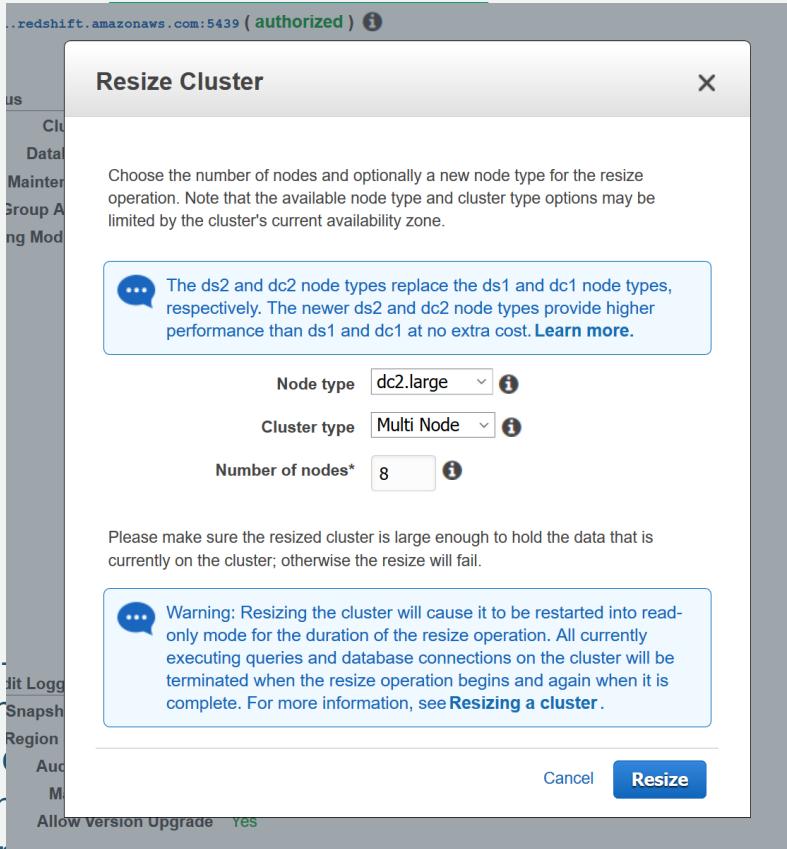
1. Cluster is put into read-only mode
2. New cluster is provisioned according to resizing needs
3. Node-to-node parallel data copy
4. Automatic SQL endpoint switchover via DNS
5. Decommission the source cluster

# 사이즈 변경



1. Cluster is put into read-only mode
2. New cluster is provisioned according to resizing needs
3. Node-to-node parallel data copy
4. Automatic SQL endpoint switchover via DNS
5. Decommission the source cluster

# 사이즈 변경



1. Cluster is put into read-only mode.
2. New cluster is provisioned.
3. Node-to-node parallel copy.
4. Automatic SQL endpoint creation.
5. Decommission the source cluster.

# 제한

## Amazon Redshift Limits

Resource	Default Limit
Nodes per cluster	101
Nodes	200
Reserved Nodes	200
Snapshots	20
Parameter Groups	20
Security Groups	20
Subnet Groups	20
Subnets per Subnet Group	20
Event Subscriptions	20

## Limit Increase 요청

The screenshot shows the Amazon Redshift Dashboard interface. On the left, there is a sidebar with links: Clusters, Snapshots, Security, Parameter Groups, Reserved Nodes, Events, and Connect Client. The 'Clusters' link is highlighted with a red oval. To the right, under the heading 'Resources', it says 'You are using the following Amazon Redshift resources'. Below this, there is a list of items: Clusters (0), Snapshots (31), Manual (31), and Automated (0). At the bottom of the dashboard, there is a 'Launch Cluster' button.

**Redshift Dashboard**

- Clusters
- Snapshots
- Security
- Parameter Groups
- Reserved Nodes
- Events
- Connect Client

**Resources**

You are using the following Amazon Redshift resources

- Clusters (0)
- Snapshots (31)
- Manual (31)
- Automated (0)

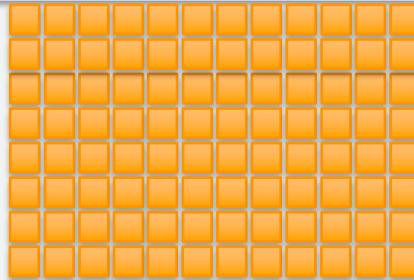
**Launch Cluster**

# 데이터 및 스키마 모델링

# Row 기반 저장소와 Column 기반 저장소

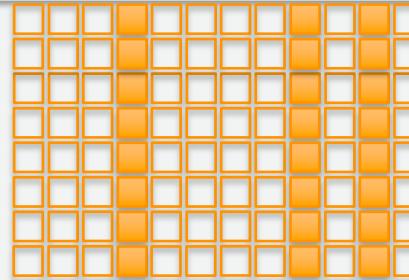
Select ID, Age, State, Amount from mytable

ID	Age	State	Amount
123	20	CA	500
345	25	WA	250
678	40	FL	125
957	37	WA	375



Row 기반 저장소

ID	Age	State	Amount
123	20	CA	500
345	25	WA	250
678	40	FL	125
957	37	WA	375

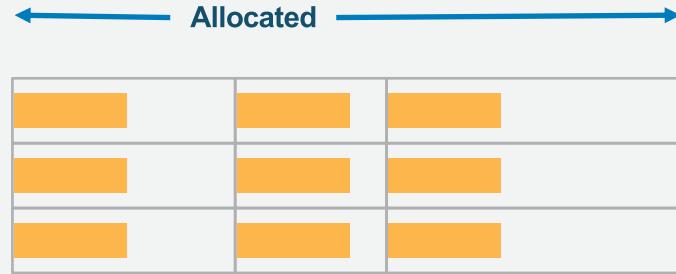


Column 기반 저장소

# Column 기반의 스토리지

Column간의 간격 최소화를 통한 성능 향상

- During queries and ingestion, the system allocates buffers based on declared column width
- Wider than needed columns mean memory is wasted
- Fewer rows fit into memory; increased likelihood of queries spilling to disk
- SVV\_TABLE\_INFO 테이블의 'max\_varchar'을 통해 가장 큰 열을 확인



# 지원 자료형

Data Type	Aliases	Description
<b>SMALLINT</b>	INT2	Signed two-byte integer
<b>INTEGER</b>	INT, INT4	Signed four-byte integer
<b>BIGINT</b>	INT8	Signed eight-byte integer
<b>DECIMAL</b>	NUMERIC	Exact numeric of selectable precision
<b>REAL</b>	FLOAT4	Single precision floating-point number
<b>DOUBLE PRECISION</b>	FLOAT8, FLOAT	Double precision floating-point number
<b>BOOLEAN</b>	BOOL	Logical Boolean (true/false)
<b>CHAR</b>	CHARACTER, NCHAR, BPCHAR	Fixed-length character string
<b>VARCHAR</b>	CHARACTER VARYING, NVARCHAR, TEXT	Variable-length character string with a user-defined limit
<b>DATE</b>		Calendar date (year, month, day)
<b>TIMESTAMP</b>	TIMESTAMP WITHOUT TIME ZONE	Date and time (without time zone)
<b>TIMESTAMPTZ</b>	TIMESTAMP WITH TIME ZONE	Date and Time (with time zone)

# JSON 지원

- Redshift tables can include arbitrary string data stored as JSON
- JSON Object or Array
- DB Functions for working with JSON
  - [IS\\_VALID\\_JSON](#)
  - [IS\\_VALID\\_JSON\\_ARRAY](#)
  - [JSON\\_ARRAY\\_LENGTH](#)
  - [JSON\\_EXTRACT\\_ARRAY\\_ELEMENT\\_TEXT](#)
  - [JSON\\_EXTRACT\\_PATH\\_TEXT](#)

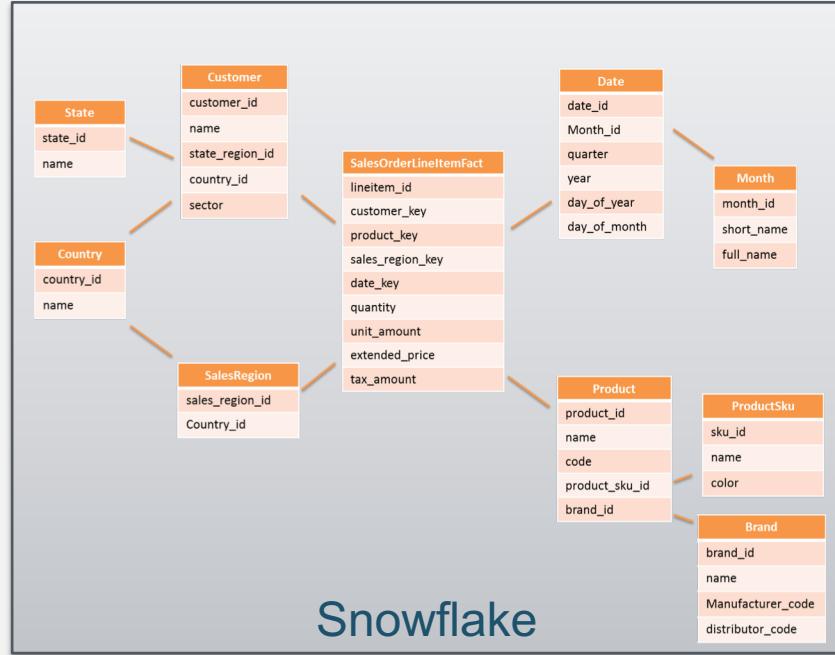
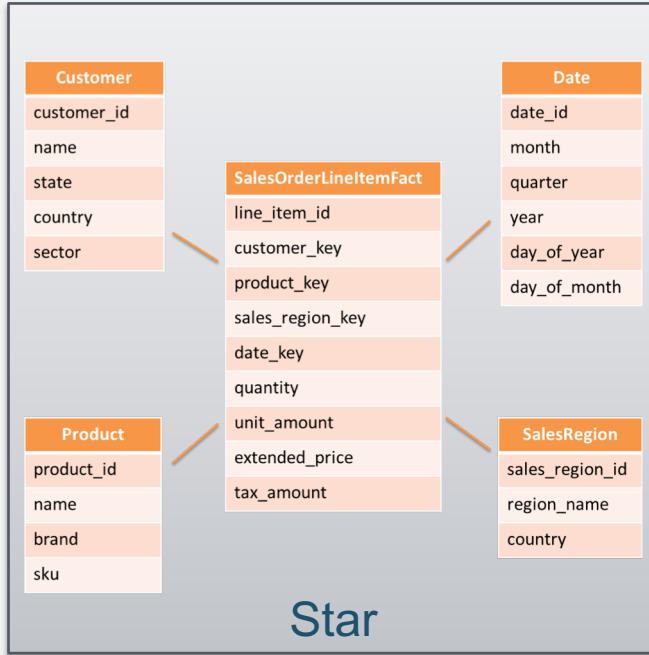
MY\_JSON\_COLUMN VARCHAR(255) NULL

```
{  
    "attribute1": "value1",  
    "attribute2": "value2",  
    "attribute3": ["value3.1", "value3.2"]  
}
```

select  
`json_extract_path_text(my_json_column, 'attribute3', 1)`

“value3.2”

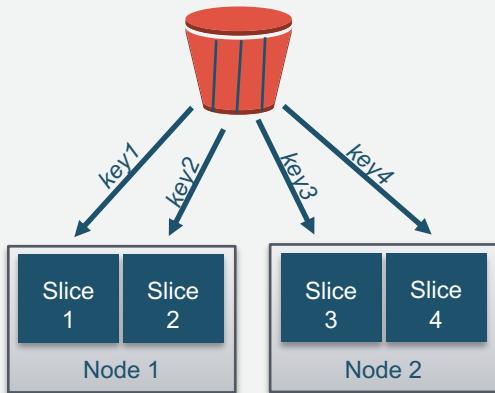
# 스키마 디자인



# Table Distribution Styles

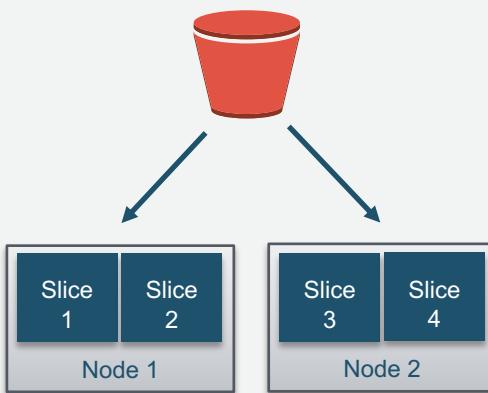
## Distribution Key

*Same key to same location*



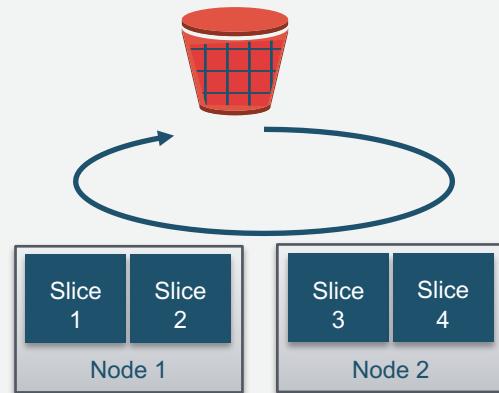
## All

*All data on every node*



## Even

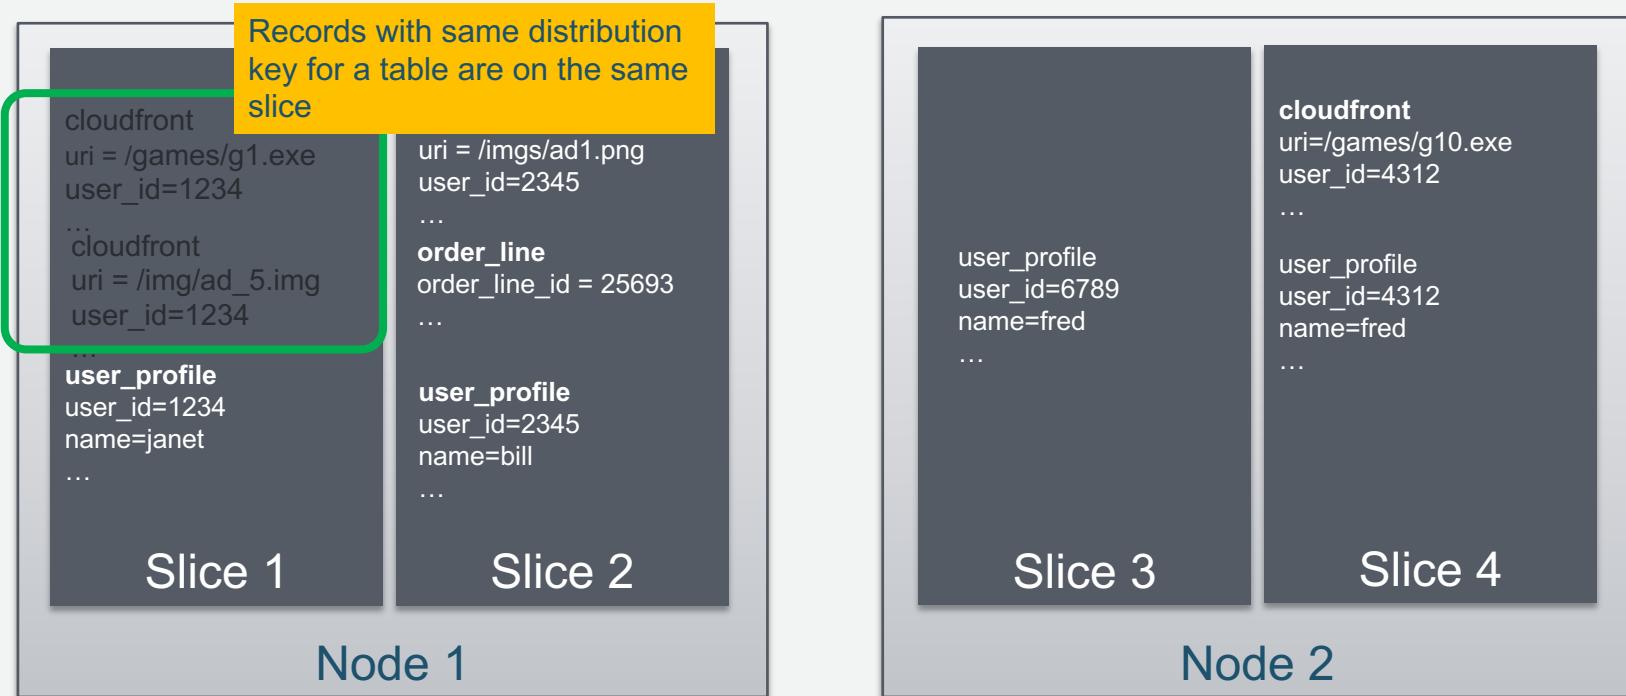
*Round robin distribution*



# Data Distribution with Distribution Keys



# Data Distribution and Distribution Keys



**Distribution Keys determine which data resides on which slices**

# Data Distribution and Distribution Keys



***Distribution Keys help with data locality for join evaluation***

# Data Distribution - Comparison

## 적절한 분산 처리

```
XN Merge (cost=1000634197251.69..1000634197251.69 rows=1 width=33)
  Merge Key: lineitem.l_shipmode
    -> XN Network (cost=1000634197251.69..1000634197251.69 rows=1 width=33)
      Send to leader
        -> XN Sort (cost=1000634197251.69..1000634197251.69 rows=1 width=33)
          Sort Key: lineitem.l_shipmode
            -> XN HashAggregate (cost=634197251.66..634197251.67 rows=1 width=33)
              -> XN Hash Join DS_DIST_NONE (cost=153674846.86..633973363.89 r
```

## DS\_DIST\_NONE, DS\_DIST\_ALL\_NONE

## Broadcast 발생

```
XN Merge (cost=9055884343615.08..9055884343615.09 rows=1 width=33)
  Merge Key: lineitemmodist.l_shipmode
    -> XN Network (cost=9055884343615.08..9055884343615.09 rows=1 width=33)
      Send to leader
        -> XN Sort (cost=9055884343615.08..9055884343615.09 rows=1 width=33)
          Sort Key: lineitemmodist.l_shipmode
            -> XN HashAggregate (cost=8055884343615.06..8055884343615.08 rows=1 width=33)
              -> XN Hash Join DS_BCAST_INNER (cost=153663142.28..8055884154841.01 rows=25
```

## DS\_BCAST\_INNER, DS\_DIST\_BOTH

# Network Broadcast

During join steps, it may be necessary for a slice to work with data that is not stored locally

Network transmission is by far the most expensive operation a query will do

The DS\_DIST\* information in an explain plan tells you how data is being accessed between tables and slices

DS_DIST*	INFO
DS_DIST_NONE	Preferred, no data transfer between nodes.
DS_DIST_ALL_NONE	One table has DISTSTYLE ALL, no data transfer between nodes.
DS_DIST_ALL_INNER	Inner table is being sent to Single Node. Very Slow.
DS_DIST_INNER	Inner table is being redistributed in an outer join.
DS_DIST_OUTER	Outer table is being redistributed in an outer join.
DS_BCAST_INNER	Inner table is being broadcast to all nodes.
DS_BCAST_BOTH	Both tables are being broadcast to all nodes.

# Distribution Style의 선택

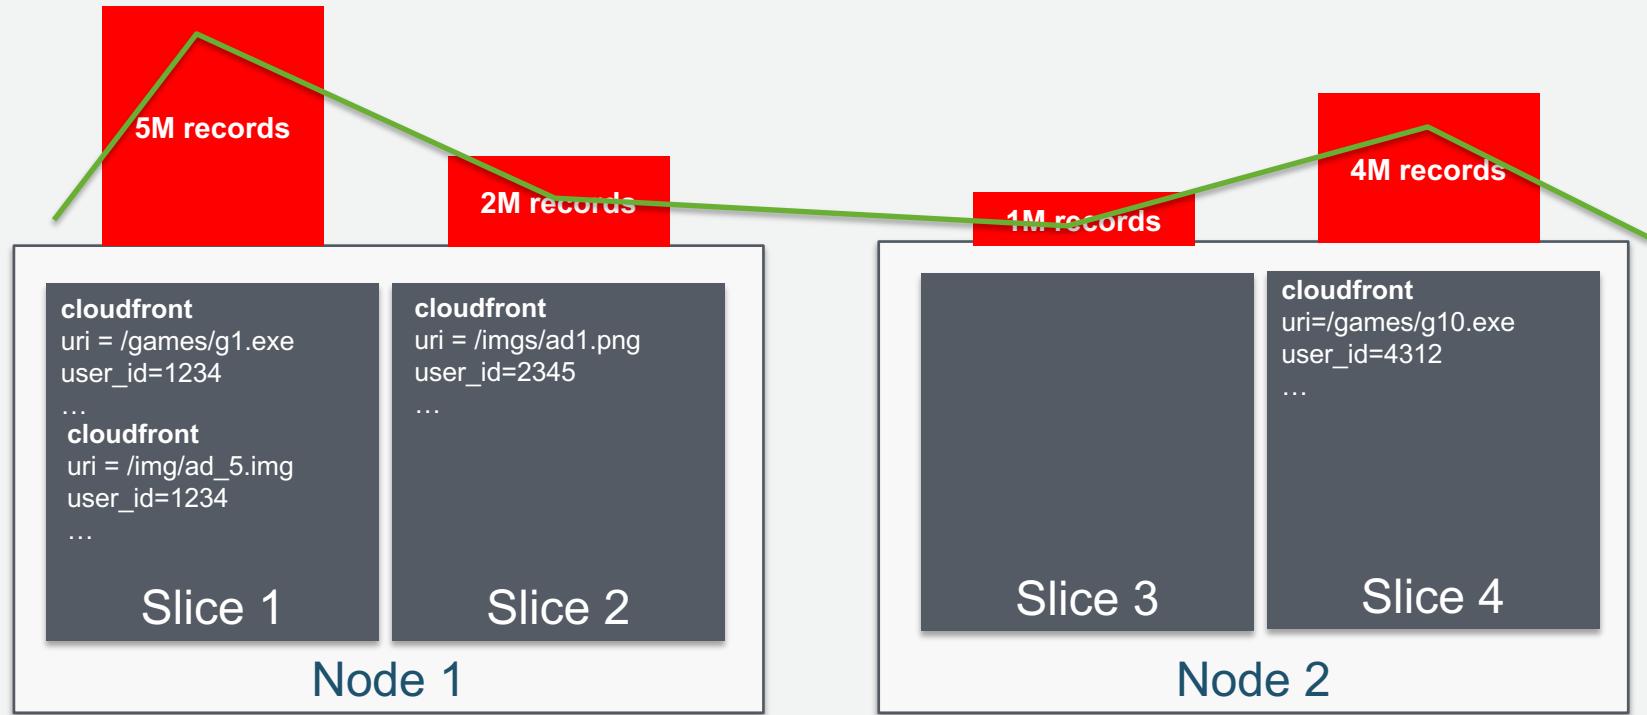
- **KEY**
  - 스타 스키마의 Fact Table과 같이 큰 테이블
  - 조인 또는 집계의 대상이 되는 큰 테이블
  - 변경이 잦은 테이블
  - KEY 값을 Join에서 사용하지 않더라도 성능 향상 효과가 있음
- **ALL**
  - 변경이 거의 없는 테이블
  - 사이즈가 작은 테이블 (몇억건 수준 미만)
  - 조인에 자주 사용되지만 특별한 Distribution Key 설정이 어려운 테이블
  - 상품의 카탈로그와 같이 작고 조인에 자주 활용되는 테이블에 대해 고려 가능
- **EVEN:** 조인 및 집계가 없는 경우

# Data Distribution and Distribution Keys



**Poor key choices lead to uneven distribution of records...**

# Data Distribution and Distribution Keys



***Unevenly distributed data cause processing imbalances!***

# Data Distribution and Distribution Keys



*Evenly distributed data improves query performance*

# Choosing a Good Distribution Key

## 목적

- Distribute data evenly across nodes
- Minimize data movement: Co-located Joins & Aggregates

## Best Practice

- Use the joined columns for largest commonly joined tables as key (example: fact table and large dimension table)
- Consider using Group By column as a key (GROUP BY clause)
- Never use a distribution key that causes severe data skew
- Choose a key with high cardinality; large number of discrete values

## Avoid

- **Keys used as equality filter** as your distribution key (Concentrates processing on one node)

# Dist Key 어떻게 선택할까?

1. 가장 자주 활용되는 쿼리를 확인합니다.  
예 "awslabs/amazon-redshift-utils/AdminScripts/top-queries.sql" 활용 등
2. 조인 및 Group By에 활용되며 고른 데이터 분포 및 높은 Cardinality 보유한 컬럼을 선택 합니다.
3. 그게 아니라면....

**"Fact Table은 Even으로 Dimension Table은 ALL로"**

"참고 : MPP 형태의 스키마는 테이블 내에  
모든 데이터를 보유하기 때문에 Distribution style 은 Even으로"  
"3정규화 형태의 Inmon 모델은 노드 수량은 줄이고  
노드 크기는 키우거나 Star Schema로 전환 "

# 참고,

1. 가장 큰 Dimension Table의 Primary키와 Fact Table의 Foreign키를 Dist Key로
2. 나머지 조인이 되는 조건의 Dimension Table은 Distribution ALL을 검토 한다.

# 성능

# 성능 향상

1. 보다 높은 하드웨어 스펙
2. 병렬화의 달성
3. 네트워크를 통한 트래픽의 최소화

**“스토리지 I/O의 최소화”**

# Sort Keys – Zone Maps

SELECT COUNT(\*) FROM LOGS WHERE DATE = '09-JUNE-2015'

Unsorted Table



MIN: 01-JUNE-2015  
MAX: 20-JUNE-2015



MIN: 08-JUNE-2015  
MAX: 30-JUNE-2015



MIN: 12-JUNE-2015  
MAX: 20-JUNE-2015



MIN: 02-JUNE-2015  
MAX: 25-JUNE-2015



MIN: 06-JUNE-2015  
MAX: 12-JUNE-2015

Sorted By Date



MIN: 01-JUNE-2015  
MAX: 06-JUNE-2015



MIN: 07-JUNE-2015  
MAX: 12-JUNE-2015



MIN: 13-JUNE-2015  
MAX: 18-JUNE-2015



MIN: 19-JUNE-2015  
MAX: 24-JUNE-2015



MIN: 25-JUNE-2015  
MAX: 30-JUNE-2015

## Zone Map?

각 Block에 대한 최대값 및  
최소값에 대한 정보 저장  
(RDB의 Index와 유사)

# Sort Keys – Single Column

Table is sorted by 1 column

[ SORTKEY ( date ) ]

Date	Region	Country
2-JUN-2015	Oceania	New Zealand
2-JUN-2015	Asia	Singapore
2-JUN-2015	Africa	Zaire
2-JUN-2015	Asia	Hong Kong
3-JUN-2015	Europe	Germany
3-JUN-2015	Asia	Korea

- Queries that use 1<sup>st</sup> column (i.e. *date*) as primary filter
- Can speed up joins and group bys
- Quickest to VACUUM

# Sort Keys – Compound

Table is sorted by 1<sup>st</sup> column , then 2<sup>nd</sup> column etc.

Date	Region	Country
2-JUN-2015	Oceania	New Zealand
2-JUN-2015	Asia	Singapore
2-JUN-2015	Africa	Zaire
2-JUN-2015	Asia	Hong Kong
3-JUN-2015	Europe	Germany
3-JUN-2015	Asia	Korea

- Queries that use 1<sup>st</sup> column as primary filter, then other cols
- Can speed up joins and group bys
- Slower to VACUUM

# Sort Keys – Interleaved

Equal weight is given to each column.

[ SORTKEY INTERLEAVED ( date, region, country) ]

Date	Region	Country
2-JUN-2015	Oceania	New Zealand
2-JUN-2015	Asia	Singapore
2-JUN-2015	Africa	Zaire
2-JUN-2015	Asia	Hong Kong
3-JUN-2015	Europe	Germany
3-JUN-2015	Asia	Korea

- Queries that use different columns in filter
- Queries get faster the more columns used in the filter (up to 8)
- Slowest to VACUUM

# Example – Compound Sort vs Interleaved Sort

- Product (**type**, **color**, **size**)
- Compound sort on **type**, **color**, **size**
  - Where **type** = ‘shirt’ and **color**=‘red’ and **size** = ‘xl’: *full advantage of sort*
  - Where **type** = ‘shirt’ and **color**=‘red’: *some efficiency*
  - Where **type** = ‘shirt’ and **size** = ‘xl’: *limited efficiency depending on cardinality of type column*
- Interleaved sort on **type**, **color**, **size**
  - Using any of **type**, **color** and **size** in a where clause will improve performance

# 압축

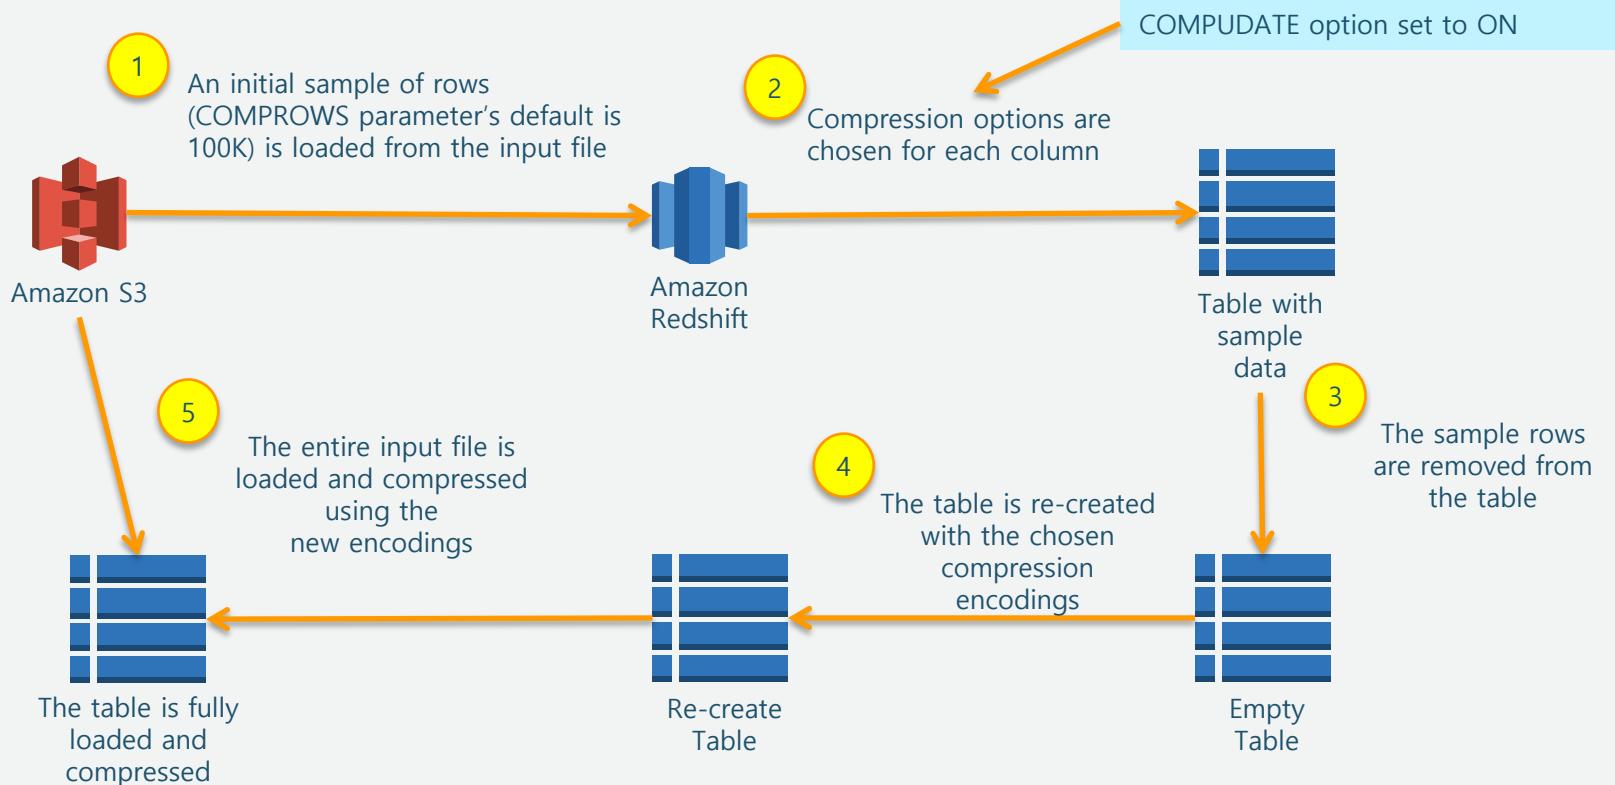
```
analyze compression listing;
```

Table	Column	Encoding
listing	listid	delta
listing	sellerid	delta32k
listing	eventid	delta32k
listing	dateid	bytedict
listing	numtickets	bytedict
listing	priceticket	delta32k
listing	totalprice	mostly32
listing	listtime	raw

- 압축을 기반으로 디스크 IO 더욱 줄일 수 있습니다.

ID	Age	State	Amount	
123	20	CA	500	
345	25	WA	250	
678	40	FL	125	
957	37	WA	375	

# Automatic Data Compression



# 쿼리 분석

EXPLAIN command followed by query string:

```
explain select avg(datediff(day, listtime, saletime)) as avgwait from sales,
listing where sales.listid = listing.listid;
          QUERY PLAN
XN Aggregate (cost=6350.30..6350.31 rows=1 width=16)
 -> XN Hash Join DS_DIST_NONE (cost=47.08..6340.89 rows=3766 width=16)
      Hash Cond: ("outer".listid = "inner".listid)
      -> XN Seq Scan on listing (cost=0.00..1924.97 rows=192497 width=12)
      -> XN Hash (cost=37.66..37.66 rows=3766 width=12)
          -> XN Seq Scan on sales (cost=0.00..37.66 rows=3766 width=12)
```

The EXPLAIN plan has the following information about query execution steps:

- Operations to be performed
- Tables and columns involved
- Data being processed

# 운영 및 모니터링

# VACUUM

ID      LOCATION\_CODE      ADDRESS

1,2,3,4	RFK,JFK,LBJ,GWB	900 Columbus,800 Washington, 700 Foxborough,600 Kansas
---------	-----------------	--

*delete from customer where id = 3;*

1,2 x	RFK,JFK,LBJ,GWB	900 Columbus,800 Washington xxxxxxxxxxxxxxxx )0 Kansas
-------	-----------------	--

*VACUUM Customer;*



1,2,4	RFK,JFK,GBW	900 Columbus,800 Washington,600 Kansas
-------	-------------	--

# 모니터링

Cluster: student1-env-rredshiftcluster Configuration Status Cluster Performance Queries Loads

Cluster: student1-env-rredshiftcluster Configuration Status Cluster Performance Queries Loads Table

Cluster Status

Time Ran

Queries

In M

Parameter Gr

Pending

Recent Eve

Filter: Last 24 Hours Search...

Terminate Query

Query Run time Start time

Time	CPU ut	Query	Run time	Start time
Mar 26 7:34		4098	35.34s	March 27, 2018 at 3:16:50 AM UTC+9
Mar 26 7:34		3913	12m 25.59s	March 27, 2018 at 2:34:18 AM UTC+9
Mar 26 7:33		2438	5.89s	March 26, 2018 at 7:56:00 PM UTC+9
		2435	1.98s	March 26, 2018 at 7:55:33 PM UTC+9
CloudWa	Network	2401	11.94s	March 26, 2018 at 7:46:02 PM UTC+9

No alarms

Query 4098 Terminate Query

Query properties

Cluster student1-env-rredshiftcluster-19z71q02xs5q9 Run time 35.34s  
Query ID 4098 Start time Tue Mar 27 03:16:50 GMT+900 2018  
Type Query End time Tue Mar 27 03:17:25 GMT+900 2018  
User admin Status Completed

SQL

```
SELECT c.c_name,
       c.c_mktsegment,
       t.prettyMonthYear,
       SUM(uv.adRevenue)
  FROM clickstream.uservisits_parquet1 AS uv
    RIGHT OUTER JOIN customer AS c
      ON c.c_custkey = uv.custKey
     INNER JOIN (
       SELECT DISTINCT d.yearnmonthnum,
                      (d.month||'-'||d.year) AS prettyMonthYear
        FROM dwdate
       WHERE d.yearnmonthnum >= 199810
     ) AS t
      ON uv.yearMonthKey = t.d_yearnmonthnum
 WHERE c.c_custkey <= 3
 GROUP BY c.c_name,
          c.c_mktsegment,
          t.prettyMonthYear,
          uv.yearMonthKey
 ORDER BY c.c_name,
          c.c_mktsegment,
          uv.yearMonthKey ASC
```

Query execution details

Plan Actual

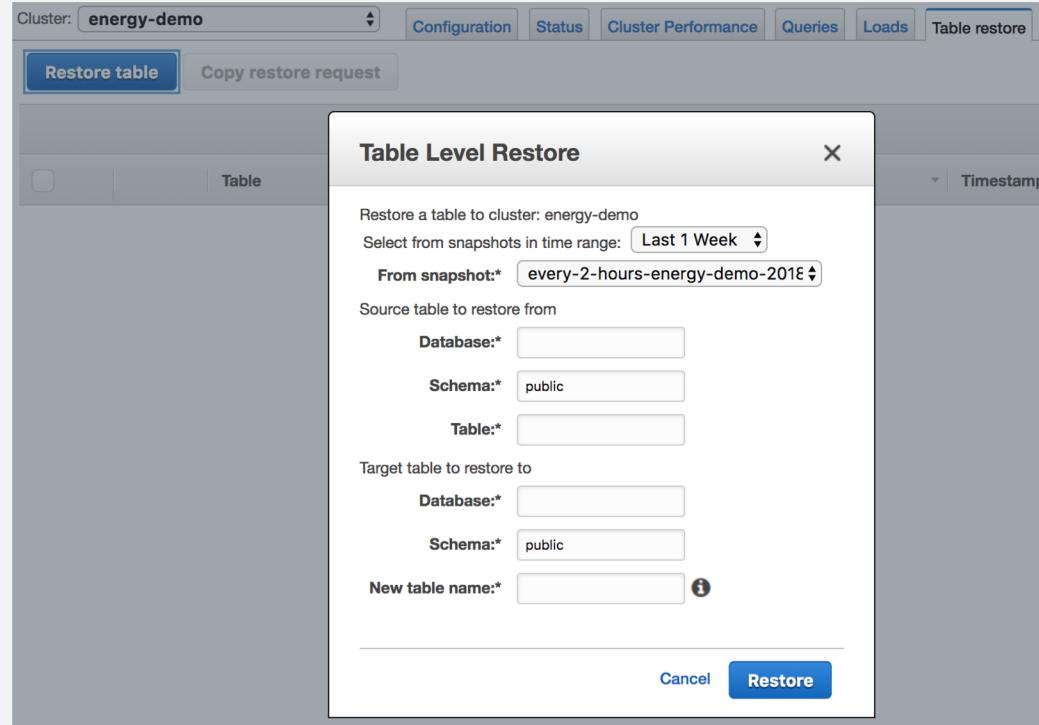
```
XN Merge (cost=1141056800867.31..1141056801917.31 rows=420000 width=78)
-> XN Network (cost=1141056800867.31..1141056801917.31 rows=420000 width=78)
  -> XN Sort (cost=1141056800867.31..1141056801917.31 rows=420000 width=78)
    -> XN HashAggregate (cost=141056760589.25..141056761639.25 rows=420000 width=78)
      -> XN Hash Join DS_DIST_ALL_NONE (cost=37.34..141056596142.86 rows=13155711 width=78)
        -> XN Hash Join DS_DIST_ALL_NONE (cost=4.50..141051766084.98 rows=375877435 width=46)
          -> XN Hash (cost=32.82..32.82 rows=7 width=36)
            -> XN Partition Loop (cost=0.00..94063327993.62 rows=3758774345000 width=16)
            -> XN Hash (cost=37.35..3.75 rows=300 width=38)
              -> XN Seq Scan PartitionInfo of clickstream.uservisits_parquet1 uv (cost=0.00..10.00)
              -> XN Seq Scan uv (cost=0.00..93969358.62 rows=3758774345 width=16)
                -> S3 Seq Scan clickstream.uservisits_parquet1 uv location="s3://kmin-redshift-spectrum-datastore-parquet1" format=PARQUET (cost=0.00..56381615.17 rows=3758774345 width=16)
                  -> XN Seq Scan on customer c (cost=0.00..3.75 rows=300 width=38)
                  -> XN Subquery Scan t (cost=0.00..32.82 rows=7 width=36)
                    -> XN Unique (cost=0.00..32.75 rows=7 width=18)
                      -> XN Seq Scan on dwdate (cost=0.00..32.43 rows=64 width=18)
```

# 단일 테이블 복원

Restore doesn't have to be full cluster

Select the Snapshot, Source Database, Schema & Table

Restore to the same cluster that created the snapshot, but a new Database, Schema & Table Name



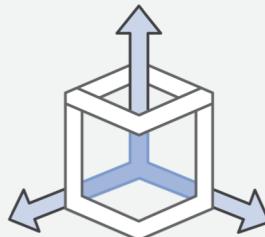
# S3 데이터레이크와 통합 Amazon Redshift Spectrum

# External Tables with Redshift Spectrum

Run SQL queries directly against data in S3 using thousands of nodes



High concurrency: Multiple clusters access same data



No ETL: Query data in-place using open file formats



Full Amazon Redshift SQL support



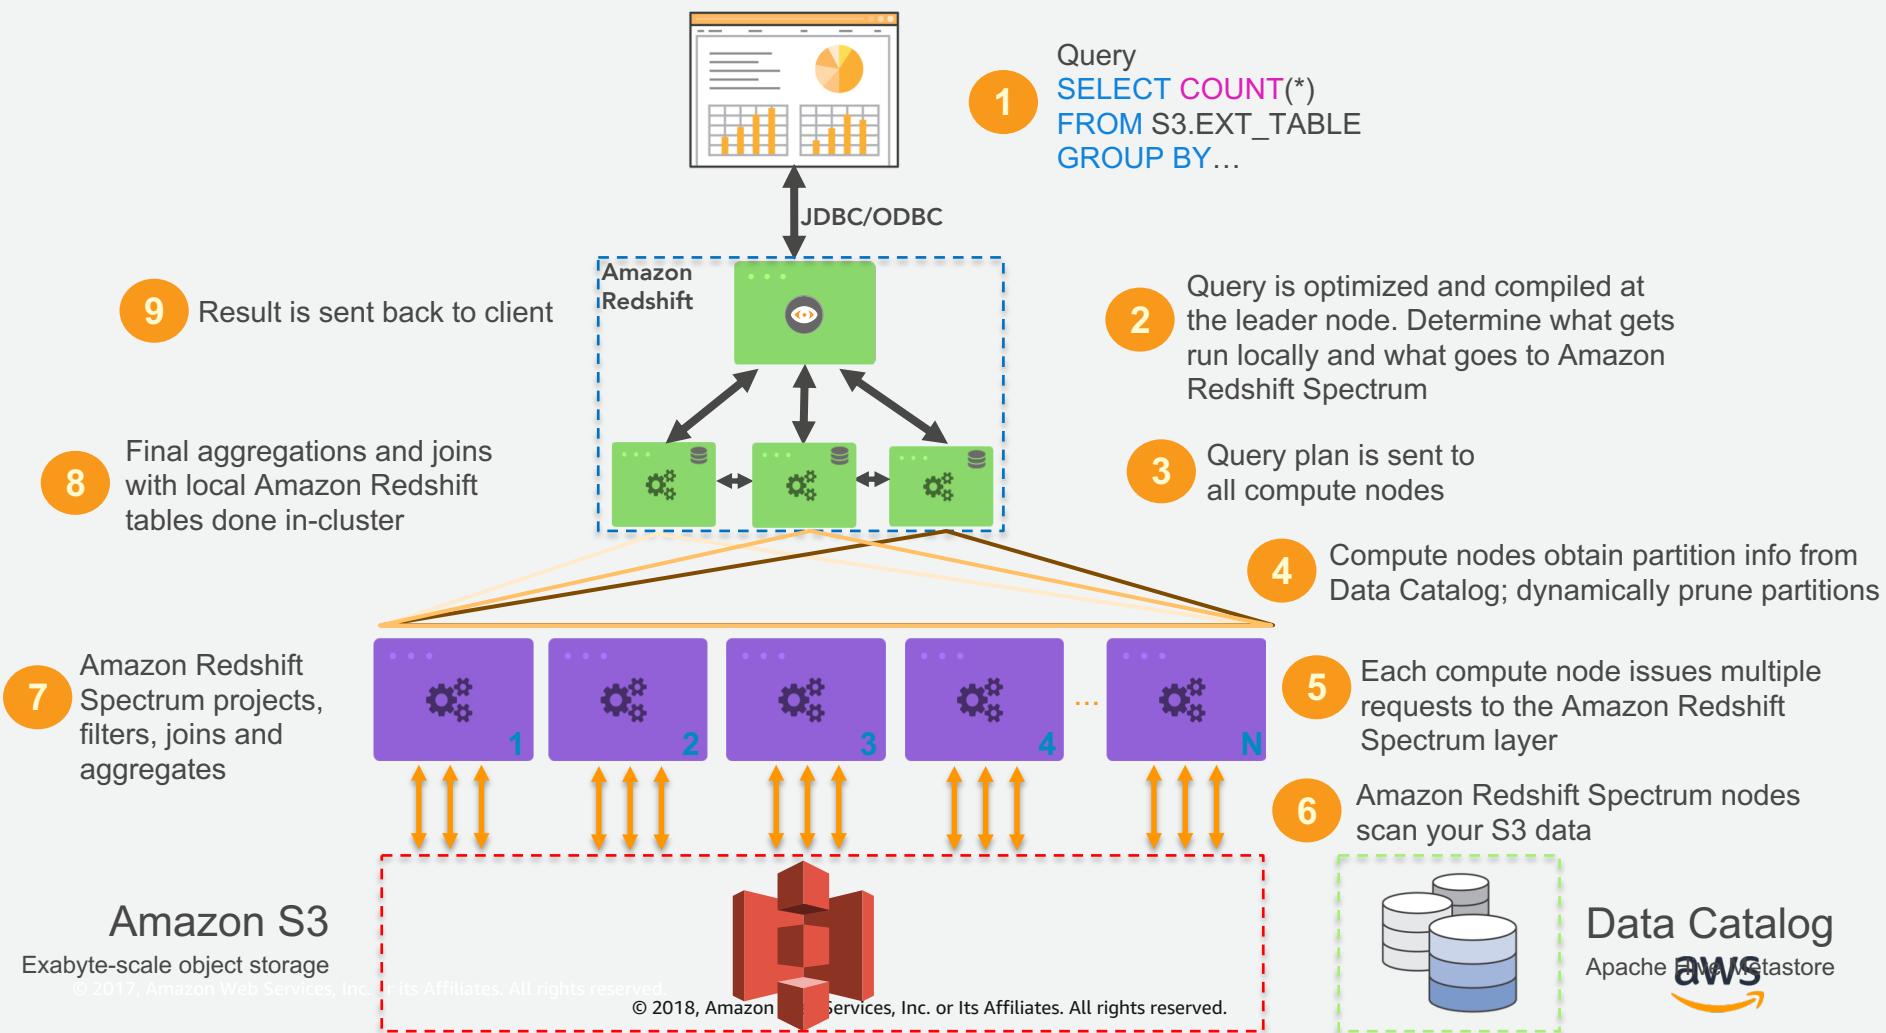
Fast @ exabyte scale



Elastic & highly available



On-demand, pay-per-query



# Redshift Spectrum에서의 성능 향상 기법

어떻게 하면 IO를 줄일것인가?

Columnar 포맷

파티셔닝

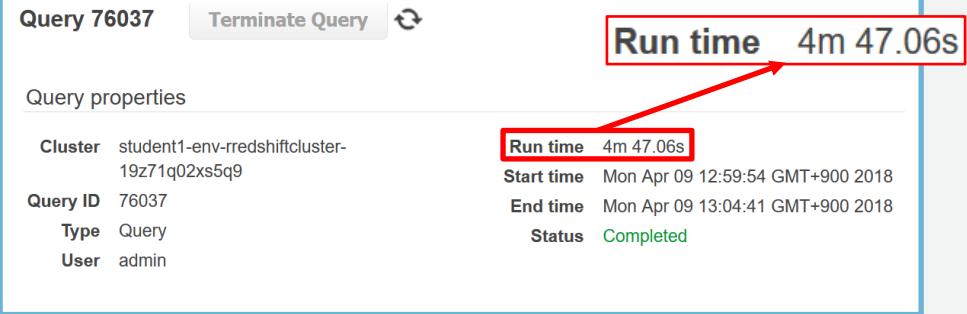
```

SELECT
    c.c_name,
    c.c_mktsegment,
    t.prettyMonthYear,
    SUM(uv.adRevenue)
FROM
    clickstream.uservisits_csv10 as uv
RIGHT OUTER JOIN
    customer as c
        ON c.c_custkey = uv.custKey
INNER JOIN
    (
        SELECT
            DISTINCT d_yearmonthnum,
            (d_month||'.'||d_year) as prettyMonthYear
        FROM
            dwdate
        WHERE
            d_yearmonthnum >= 199810
    ) as t
        ON uv.yearMonthKey = t.d_yearmonthnum
WHERE
    c.c_custkey <= 3
GROUP BY
    c.c_name,
    c.c_mktsegment,
    t.prettyMonthYear,
    uv.yearMonthKey
ORDER BY
    c.c_name,
    c.c_mktsegment,
    uv.yearMonthKey ASC

```

XN Merge (cost=1141056800867.31..1141056801917.31 rows=420000 width=78)  
Network (cost=1141056800867.31..1141056801917.31 rows=420000 width=78)  
N Sort (cost=1141056800867.31..1141056801917.31 rows=420000 width=78)  
XN HashAggregate (cost=141056760589.25..141056761639.25 rows=420000 width=78)  
-> XN Hash Join DS\_DIST\_ALL\_NONE (cost=37.34..141056596142.86 rows=13155711 width=78)  
-> XN Hash Join DS\_DIST\_ALL\_NONE (cost=4.50..14105176609.198 rows=375877435 width=46)  
-> XN Hash (cost=3.75..3.75 rows=300 width=38)  
-> XN Seq Scan on uv (cost=0.00..3.75 rows=300 width=38)  
-> XN Hash (cost=32.75..32.75 rows=7 width=18)  
-> XN Subquery Scan (cost=32.75..32.75 rows=7 width=18)  
-> XN Unique (cost=0.00..32.75 rows=7 width=18)  
-> XN Seq Scan on dwdate (cost=0.00..32.43 rows=64 width=18)  
-> XN Seq Scan PartitionInfo of clickstream.uservisits\_csv10 uv (cost=0.00..10.00 rows=1000 width=0)  
-> XN S3 Query Scan uv (cost=0.00..93969358.52 rows=3758774345 width=16)

**rows=3758774345**



- 결과 : 약 4분 47초
- CSV 포맷
- 파티셔닝 없음

```

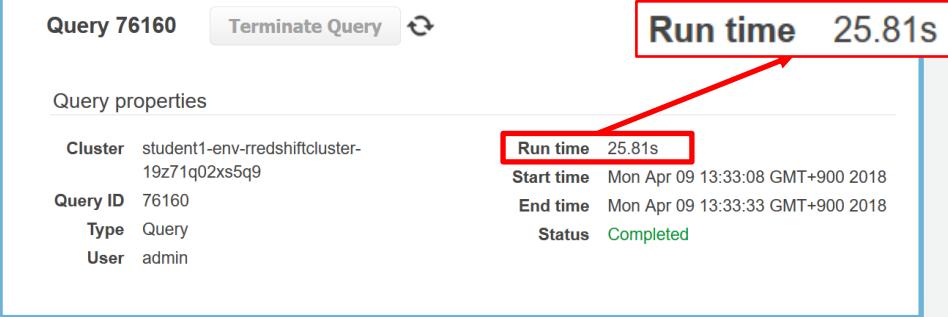
SELECT
    c.c_name,
    c.c_mktsegment,
    t.prettyMonthYear,
    SUM(uv.adRevenue)
FROM
    clickstream.uservisits_parquet1 as uv
RIGHT OUTER JOIN
    customer as c
        ON c.c_custkey = uv.custKey
INNER JOIN
    (
        SELECT
            DISTINCT d_yearmonthnum,
            (d_month||'.'||d_year) as prettyMonthYear
        FROM
            dwdate
        WHERE
            d_yearmonthnum >= 199810
    ) as t
        ON uv.yearMonthKey = t.d_yearmonthnum
WHERE
    c.c_custkey <= 3
GROUP BY
    c.c_name,
    c.c_mktsegment,
    t.prettyMonthYear,
    uv.yearMonthKey
ORDER BY
    c.c_name,
    c.c_mktsegment,
    uv.yearMonthKey ASC

```

```

XN Merge (cost=1141056800867.31..1141056801917.31 rows=420000 width=78)
    -> XN Network (cost=1141056800867.31..1141056801917.31 rows=420000 width=78)
        (cost=1141056800867.31..1141056801917.31 rows=420000 width=78)
        HashAggregate (cost=141056760589.25..141056761639.25 rows=420000 width=78)
        Hash Join DS_DIST_ALL_NONE (cost=37.34..1410565596142.86 rows=13155711 width=78)
        -> XN Hash Join DS_DIST_ALL_NONE (cost=4.50..141051766084.08 rows=375877435 width=46)
        -> XN Hash (cost=3.75..3.75 rows=300 width=38)
        -> XN Seq Scan on customer c (cost=0.00..3.75 rows=300 width=38)
        -> XN Hash (cost=3.75..3.75 rows=300 width=38)
        -> XN Subquery Scan (cost=3.75..3.75 rows=300 width=38)
        -> XN Unique (cost=0.00..0.25 rows=1 width=16)
        -> XN Seq Scan on dwdate (cost=0.00..32.43 rows=64 width=18)
        -> XN Seq Scan PartitionInfo of clickstream.uservisits_parquet1 uv (cost=0.00..10.00 rows=1000
width=0)
        -> XN S3 Query Scan uv (cost=0.00..93969358.62 rows=375877435 width=16)

```



- 결과 : 약 26초
- Parquet 포맷 (Columnar 포맷)
- 파티셔닝 없음

Overview					
<input type="text"/> Type a prefix and press Enter to search. Press ESC to clear.					
<a href="#">Upload</a>		<a href="#">+ Create folder</a>			
<u><a href="#">customer=5</a></u> / <u><a href="#">visitYearMonth=199607</a></u>					
Viewing 1 to 10					
<input type="checkbox"/> Name	Last modified	Size	Storage class		
<input type="checkbox"/> <a href="#">part-03900-64803196-a9da-49c8-ad41-cd4838b9f56a.csv</a>	Mar 20, 2018 11:36:25 AM GMT+0900	126.6 MB	Standard		
<input type="checkbox"/> <a href="#">part-03901-64803196-a9da-49c8-ad41-cd4838b9f56a.csv</a>	Mar 20, 2018 11:36:26 AM GMT+0900	126.6 MB	Standard		
<input type="checkbox"/> <a href="#">part-03902-64803196-a9da-49c8-ad41-cd4838b9f56a.csv</a>	Mar 20, 2018 11:36:27 AM GMT+0900	126.6 MB	Standard		
<input type="checkbox"/> <a href="#">part-03903-64803196-a9da-49c8-ad41-cd4838b9f56a.csv</a>	Mar 20, 2018 11:36:29 AM GMT+0900	126.6 MB	Standard		
<input type="checkbox"/> <a href="#">part-03904-64803196-a9da-49c8-ad41-cd4838b9f56a.csv</a>	Mar 20, 2018 11:36:30 AM GMT+0900	126.6 MB	Standard		
<input type="checkbox"/> <a href="#">part-03905-64803196-a9da-49c8-ad41-cd4838b9f56a.csv</a>	Mar 20, 2018 11:36:31 AM GMT+0900	93.7 MB	Standard		
<input type="checkbox"/> <a href="#">part-03906-64803196-a9da-49c8-ad41-cd4838b9f56a.csv</a>	Mar 20, 2018 11:36:32 AM GMT+0900	101.4 MB	Standard		

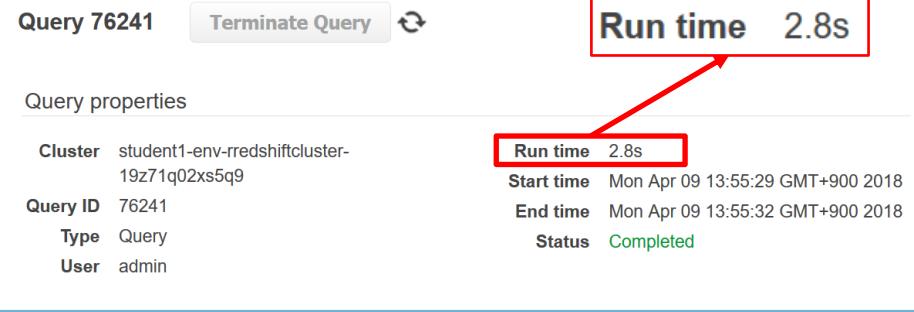
```

SELECT
    c.c_name,
    c.c_mktsegment,
    t.prettyMonthYear,
    uv.totalRevenue
FROM
    ((SELECT
        customer,
        visitYearMonth,
        SUM(adRevenue) AS totalRevenue
    FROM
        clickstream.uservisits_parquet1
    WHERE
        customer <= 3
        AND visitYearMonth >= 199810
    GROUP BY
        customer,
        visitYearMonth) AS uv
    RIGHT OUTER JOIN
        customer AS c
        ON c.c_custkey = uv.customer
    INNER JOIN
        (
            SELECT
                DISTINCT d_yeарmonthnum,
                (d_month || '-' || d_year) AS prettyMonthYear
            FROM
                dwdate
            WHERE
                d_yeарmonthnum >= 199810
        ) AS t
        ON uv.visitYearMonth = t.d_yeарmonthnum
    )
ORDER BY
    c.c_name,
    c.c_mktsegment,
    uv.visitYearMonth ASC

```

XN Merge (cost=1000094008880.16..1000094008880.18 rows=7 width=78)  
 -> XN Network (cost=1000094008880.16..1000094008880.18 rows=7 width=78)  
 -> XN Sort (cost=1000094008880.16..1000094008880.18 rows=7 width=78)  
 -> XN Hash Join DS\_DIST\_ALL\_NONE (cost=94008878.97..94008880.06 rows=7 width=78)  
 -> XN Hash Join DS\_DIST\_ALL\_NONE (cost=93971378.97..93971379.61 rows=7 width=48)  
 -> XN Subquery Scan uv (cost=93971346.13..93971346.42 rows=23 width=16)  
 -> XN HashAggregate (cost=93971346.13..93971346.19 rows=1 width=78)  
 -> XN Hash (cost=32.82..32.82 rows=7 width=36)  
 -> XN Subquery Scan t (cost=0.00..32.82 rows=7 width=36)  
 -> XN Unique (cost=0.00..32.75 rows=7 width=18)  
 -> XN Seq Scan on dwdate (cost=0.00..32.43 rows=64 width=18)  
 -> XN Hash (cost=30000.00..30000.00 rows=3000000 width=30)  
 -> XN Seq Scan on customer c (cost=0.00..30000.00 rows=3000000 width=38)  
 -> XN Seq Scan PartitionInfo of clickstream.uservisits\_parquet1 (cost=0.00..17.50 rows=112 width=8)  
 -> XN S3 Query Scan uservisits\_parquet1 (cost=46984679.32..46984689.32 rows=1000 width=8)

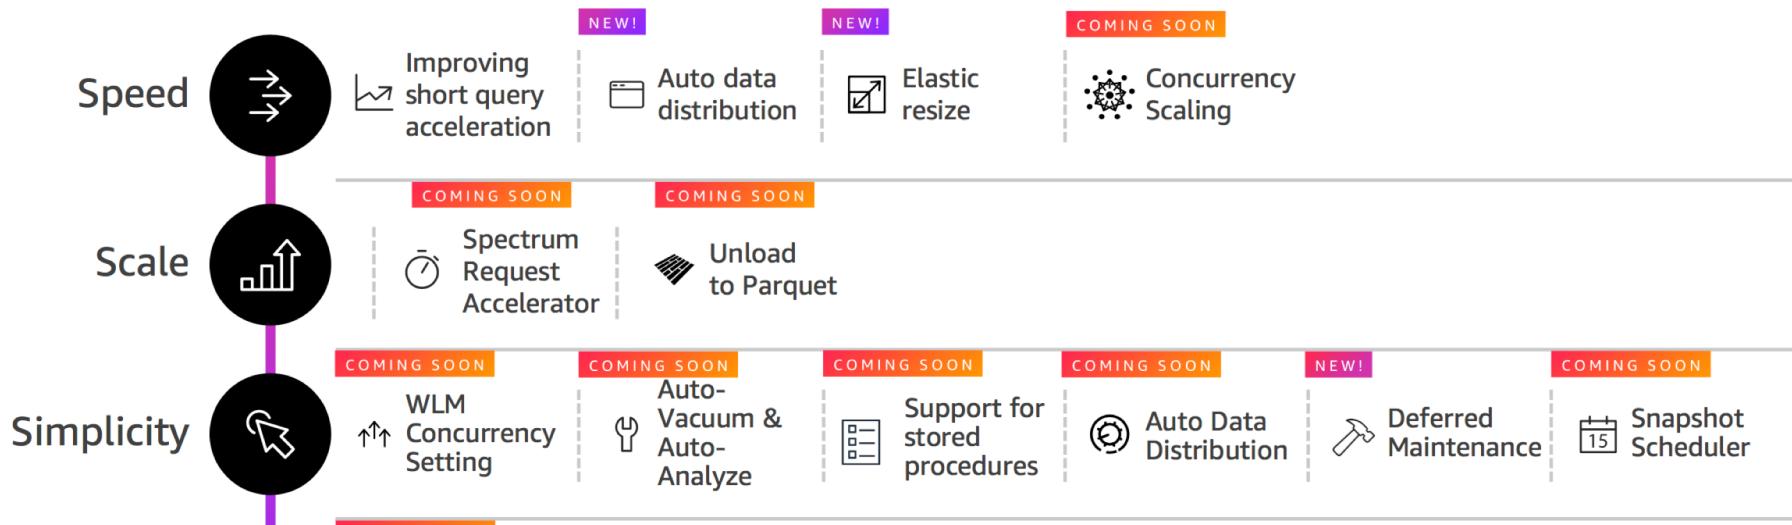
**rows=3000000**



- 결과 : 약 3초
- Parquet 포맷 (Columnar 포맷)
- 파티셔닝 : customer, visitYearMonth

# Amazon Redshift

## New features

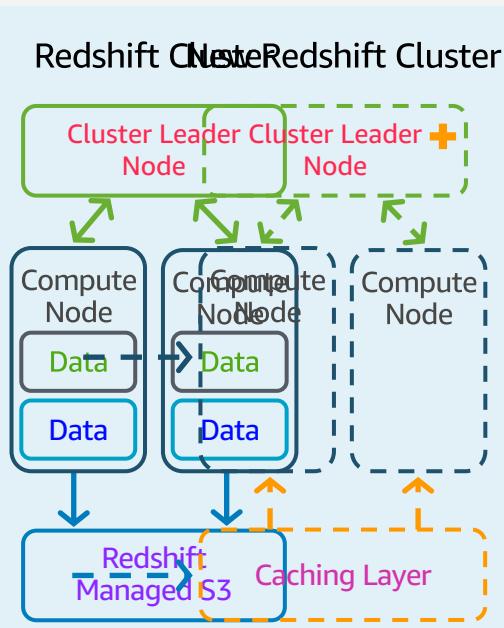


NEW!



# Amazon Redshift Concurrency Scaling (Preview)

Consistently fast performance at virtually unlimited concurrency



Automatically spins up additional clusters on-demand

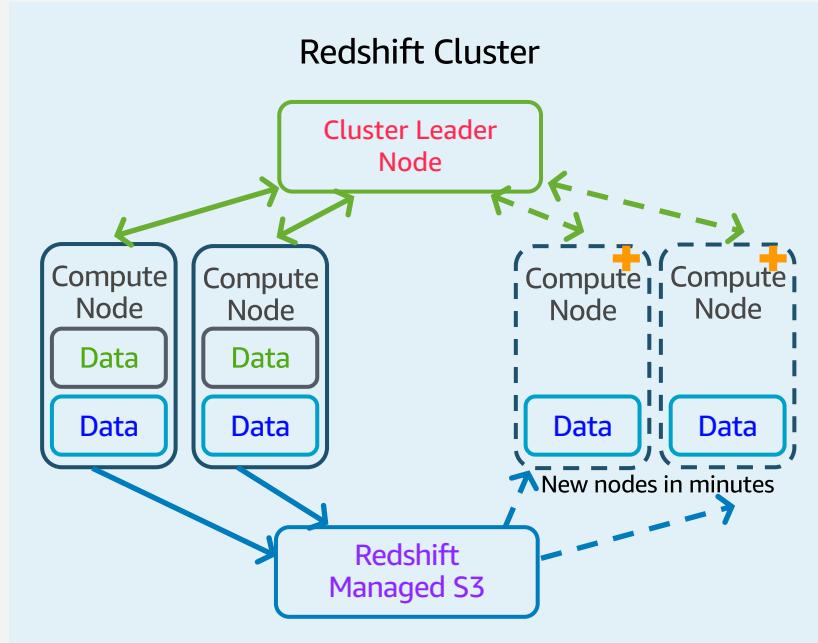
Handles virtually unlimited number of concurrent users

Accrued minutes make it free for most customers



# Amazon Redshift Elastic Resize (GA) NEW!

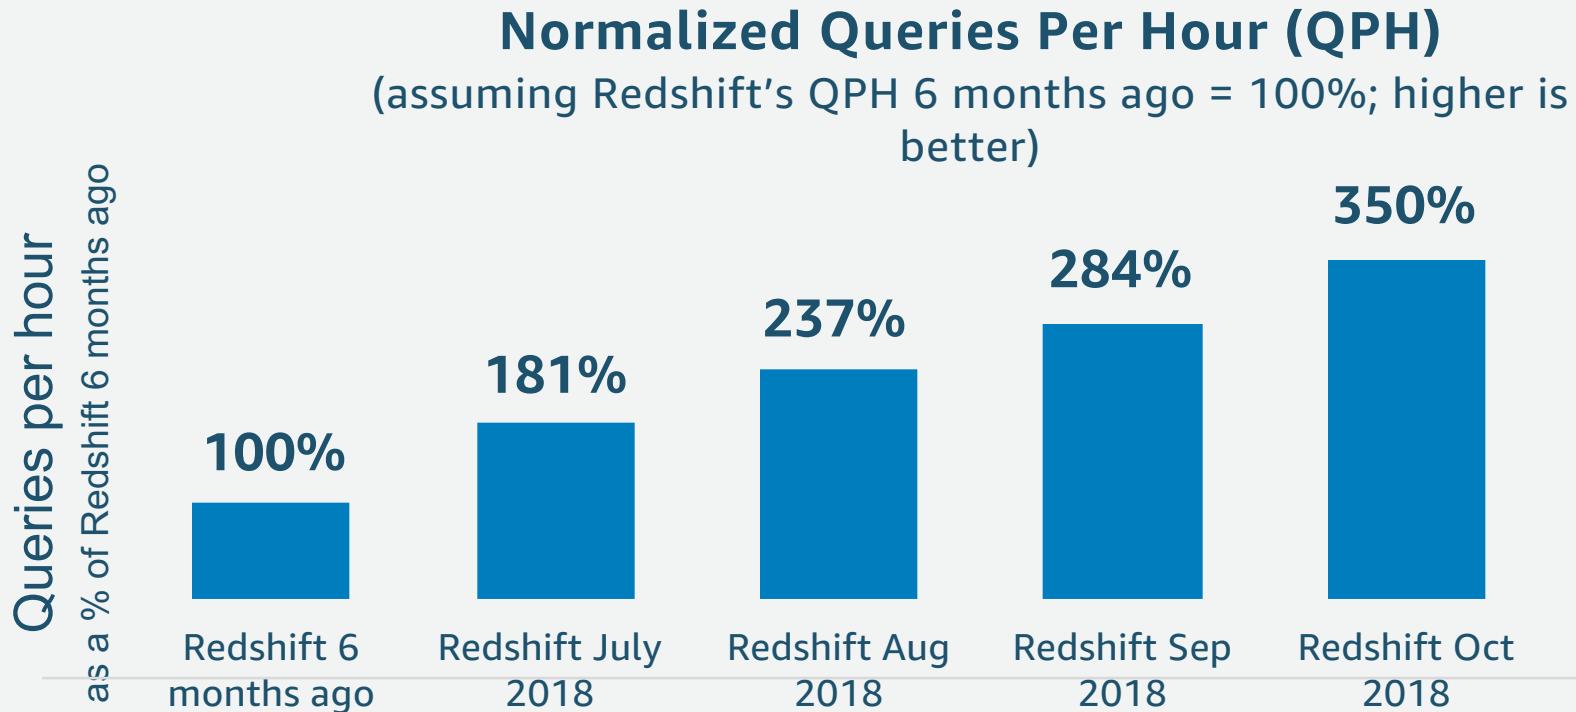
Quickly scale up or down to increase performance on-demand



Add/remove additional nodes to cluster in minutes

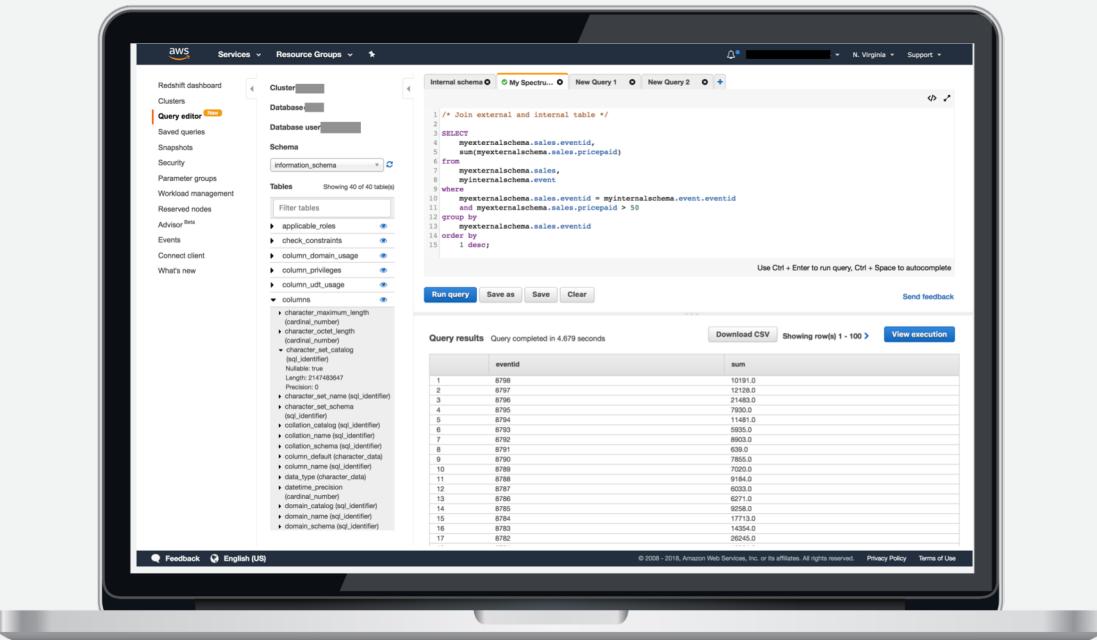
Available today

# Amazon Redshift is >3x faster than 6 months ago



# Redshift Query Editor

Launched in October!



**Query data  
directly from  
the AWS console**

Results are instantly  
visible within the console

No need to install  
and setup an external  
JDBC/ODBC client

# Amazon Redshift intelligent administration



Automates data distribution in tables for improved performance and disk space utilization.

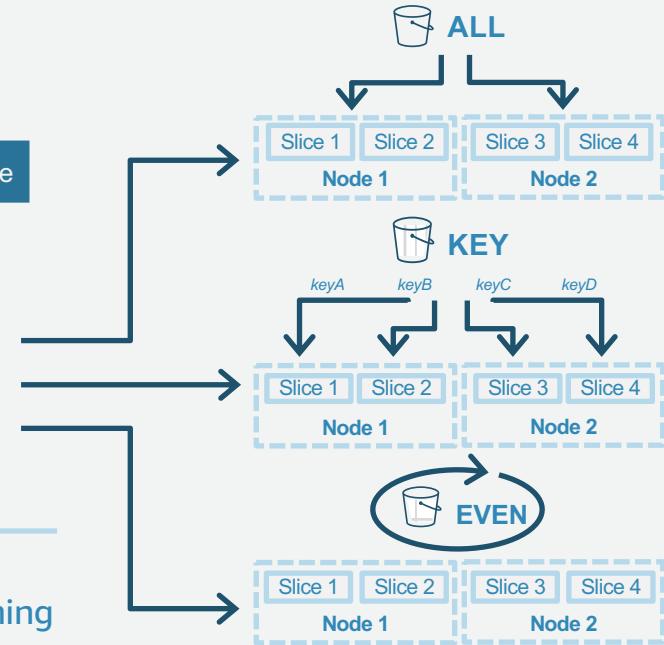
No more messing with distkeys!

Coming Soon!

 **Amazon Redshift**  
recommended distribution key

Advise

Provides intelligent recommendations for tuning based on continuous workload analysis.



# Amazon Redshift intelligent maintenance

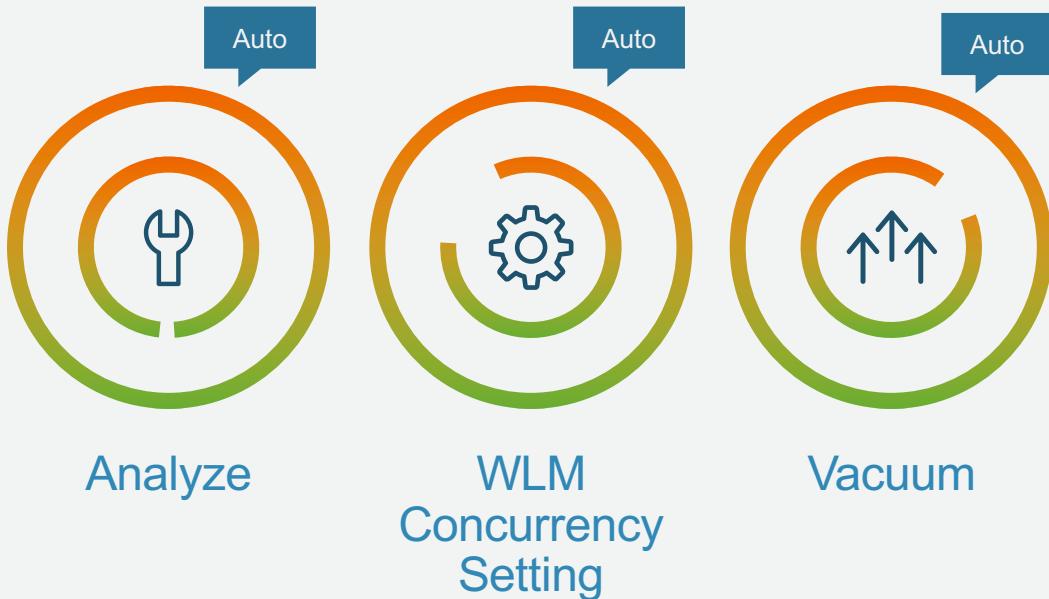
Coming Soon!



Maintenance processes like vacuum and analyze will automatically run in the background.

Moving towards zero-maintenance.

Amazon Redshift will automatically adjust the WLM concurrency setting to deliver optimal throughput.



# Run stored procedures in Amazon Redshift



Bring your existing **Stored Procedure** and run in Redshift.

Migrating to Redshift  
is even easier!

Amazon Redshift will support Stored Procedure in PL/pgSQL format, enabling you to bring your existing Stored Procedure to Amazon Redshift.

Coming Soon!

**Support for stored procedure provides the ability to run code** where the data is to efficiently run ETL, data validation, and custom business logic.



# Amazon Redshift Spectrum

Extend the data warehouse to exabytes of data in S3 Data Lake

No loading required

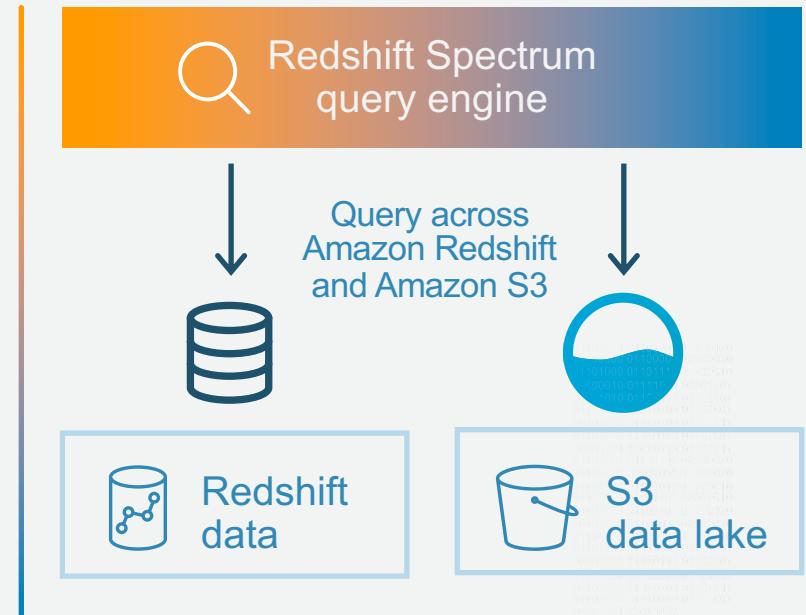
Scale compute and storage separately

Directly query data stored in Amazon S3

Parquet, ORC, Avro, Grok, and CSV data formats

- Unload to Parquet
- Spectrum Request Accelerator

Coming Soon!



# Security is built-in



End-to-end encryption



Integration with AWS Key Management Service

## Select compliance certifications\*



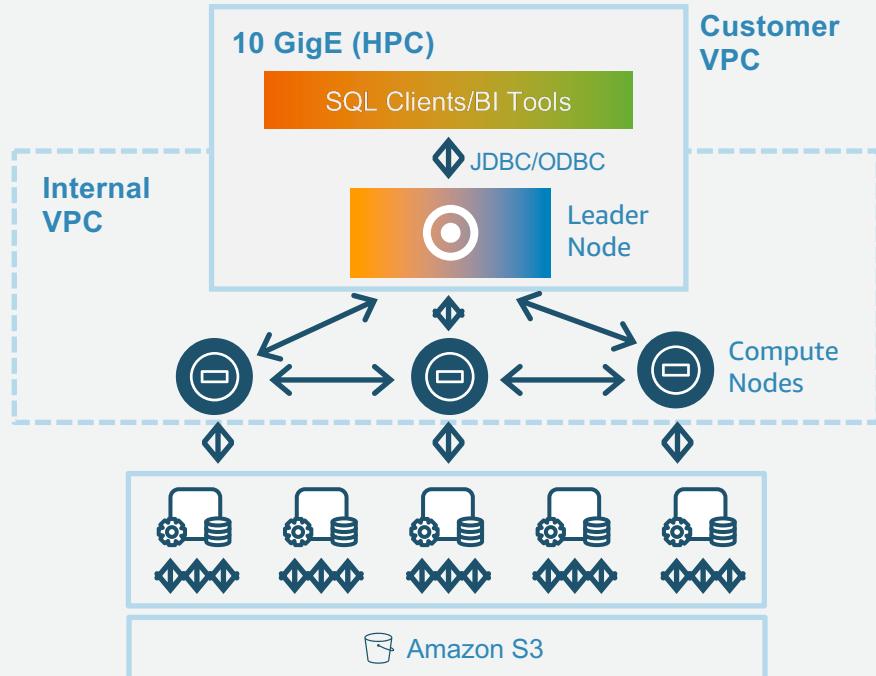
FedRAMP



HIPAA  
COMPLIANT



## Network Isolation



\*Full list of compliance certifications available here: <https://aws.amazon.com/compliance>

# 참고 URLs

- Amazon Redshift 를 위한 10가지 성능 튜닝 기법

<https://aws.amazon.com/ko/blogs/korea/top-10-performance-tuning-techniques-for-amazon-redshift/>

- Redshift 운영 시 참고 스크립트

<https://github.com/awslabs/amazon-redshift-utils>

